

The Case for CatBoost: An Adversarial Evaluation for High-Cardinality Spatio-Temporal Classification Under NTRO Constraints

Problem Deconstruction: A Multi-Layered Challenge

The selection of an optimal machine learning model for classifying NASA FIRMS thermal hotspots is not a standard classification task; it is a multi-layered challenge defined by the unique properties of the data and the stringent operational constraints of the National Technical Research Organization (NTRO). The problem space is characterized by three primary dimensions: the inherent complexity of high-cardinality spatial tabular data, the practical limitations of the computational environment, and the non-negotiable requirements for model transparency and trustworthiness. Understanding these dimensions is critical to navigating the landscape of potential models and arriving at a scientifically valid and operationally feasible solution.

The first and most significant challenge stems from the data's structure. The aggregation of VIIRS thermal hotspots onto an H3 hexagonal grid system creates a high-cardinality categorical feature on a massive scale ¹⁶. An H3 grid can partition the Earth's surface into millions of discrete hexagonal cells, each identified by a unique alphanumeric code ¹¹⁴⁴³². When these H3 cell identifiers are treated as a categorical feature for a machine learning model, they represent a feature with an extremely large number of possible values. This immediately presents a formidable obstacle for many common modeling techniques. Standard methods for handling categorical variables, such as One-Hot Encoding (OHE), become computationally intractable. OHE transforms a single categorical feature into a large number of binary columns, one for each unique category. For millions of H3 cells, this would result in a feature matrix of prohibitive size, guaranteed to cause Out-of-Memory (OOM) errors in the constrained environment of a Kaggle notebook, which operates under strict CPU and RAM limitations ²⁹⁸²⁹⁹. Any viable model must therefore possess an intrinsic ability to process such high-cardinality categorical data efficiently, without requiring memory-intensive preprocessing steps that distort the underlying information.

The second major challenge is the nature of the data itself: spatio-temporal autocorrelation. Thermal hotspots are not randomly distributed across space and time; they exhibit strong clustering patterns. Geographically, persistent industrial flares, clusters of agricultural burning, and wildfire fronts create dense spatial patterns ^{1 5}. Temporally, a hotspot detected today has a high probability of being detected tomorrow if its source is stable, such as a gas flare or an active volcano ^{53 135}. This inherent correlation between observations violates a fundamental assumption of many statistical and machine learning models—that data points are independent and identically distributed. A naive model, particularly a powerful ensemble method like gradient boosting, can easily overfit to these spatial and temporal patterns in the training data ⁸. Instead of learning the generalizable rules that distinguish a Gas Flare from an Agricultural burn based on their features (e.g., `frp`, `persistence_90d`, `bright_ti4`), the model may simply memorize the geographic locations present in the training set. This leads to a dangerous form of target leakage, where information about the target variable from one location inadvertently influences predictions at a nearby, but unseen, location ⁴³⁰⁴⁸⁹. The result is a model that achieves artificially inflated performance metrics on the test set (if it is split randomly) but fails catastrophically when deployed to classify hotspots in new geographic areas, as it has not learned the underlying physical phenomena but rather the layout of the training map ¹¹⁶. Therefore, any selected model must incorporate mechanisms to explicitly mitigate this spatio-temporal leakage to ensure true generalization.

The third dimension of the challenge is defined by the strict operational and ethical constraints imposed by the NTRO. The requirement for "defense-grade explainability" fundamentally precludes the use of black-box models, such as complex neural networks or deep ensembles whose internal decision-making processes are opaque ^{21 239}. In a national security context, a model's prediction cannot be accepted on faith alone; an intelligence analyst must be able to audit and verify the reasoning behind a classification ¹⁷³²⁹³. This necessitates a model that either has intrinsic interpretability or can be reliably explained using a post-hoc method that is both mathematically sound and analytically verifiable. The Shapley Additive Explanations (SHAP) framework is frequently cited as a state-of-the-art method for providing such explanations, as it assigns each feature an importance value for a particular prediction based on game theory, ensuring consistency and local accuracy ¹⁸⁴²⁹¹. The output must translate into clear, logical statements that an analyst can understand and validate, such as "This hotspot was classified as a 'Wildfire' primarily because its Fire Radiative Power (FRP) is in the top decile of its persistence window and its brightness temperature (`bright_ti4`) is below 300K." This level of forensic detail is essential for building trust and enabling human-in-the-loop decision-making ¹⁶⁸³⁴⁵. Furthermore, the entire project is bounded by a 3-week timeline for a

single ML engineer. This constraint makes overly complex solutions impractical. Architectures like Graph Neural Networks (GNNs) or custom-built spatio-temporal transformers, while theoretically powerful, involve significant engineering overhead, require specialized libraries, and are highly unlikely to converge or perform well within the CPU-only Kaggle environment [23](#) [83](#) [122](#). The ideal solution must strike a balance, offering state-of-the-art performance on this specific data type while remaining pragmatic and implementable within the given timeframe.

In summary, the problem is not merely to find a model that classifies hotspots accurately. It is to find a model that can gracefully handle millions of unique spatial categories without consuming excessive memory, that can learn generalizable rules despite strong spatial and temporal autocorrelation, and that can produce transparent, verifiable, and trustworthy explanations for its decisions. These interconnected challenges define the search criteria and elevate the task beyond simple benchmarking to a rigorous adversarial evaluation of candidate algorithms against a demanding set of real-world constraints.

CatBoost: The Algorithmic Solution to High-Cardinality Spatial Data

The selection of a suitable machine learning model for this problem hinges on finding an algorithm capable of addressing the core challenges of high-cardinality categorical features and spatio-temporal autocorrelation without compromising on feasibility or explainability. While numerous Gradient Boosted Decision Tree (GBDT) libraries exist, CatBoost stands out due to two foundational innovations detailed in its seminal research paper: Ordered Target Statistics for efficient categorical encoding and Ordered Boosting for preventing target leakage [201205](#). These features are not incremental improvements but are specifically engineered to solve the exact problems presented by the FIRMS dataset aggregated on H3 grids. Peer-reviewed literature consistently validates CatBoost's superiority in scenarios involving heterogeneous data, which aligns perfectly with the mixed-type feature set of this classification task [180](#).

The primary advantage of CatBoost lies in its native handling of categorical features through Target Statistics (TS), often referred to as target encoding [215](#). Unlike XGBoost or LightGBM, which typically require external preprocessing to convert categorical variables into numerical representations, CatBoost automates this process internally during the training loop [85](#) [218](#). For a high-cardinality feature like an H3 hexagon index, TS works by

replacing each categorical value with the mean of the target variable for that specific category, calculated on a subset of the data used for training previous trees ⁴⁸⁸. This approach offers a profound advantage over One-Hot Encoding (OHE). OHE would generate a separate binary column for every unique H3 cell, leading to a sparse and astronomically wide feature matrix that would crash any CPU-based environment ²⁹⁸. TS, conversely, transforms each category into a single continuous numerical feature representing its average predictive power. This drastically reduces the memory footprint and avoids the combinatorial explosion of features associated with OHE ³⁵⁴. Critically, CatBoost further refines this process with its ordered version of TS. To prevent target leakage at the most basic level, it calculates the target statistics for a given data point using only the preceding data points in a random permutation of the dataset ³²⁶³²⁷. This ensures that a sample's own target value does not influence its encoded representation, a subtle but crucial step in building a robust model. By natively supporting this memory-efficient and statistically sound encoding method, CatBoost sidesteps the need for cumbersome and potentially misleading manual preprocessing, saving significant development time and mitigating the risk of error within the tight 3-week timeline.

The second pillar of CatBoost's design is Ordered Boosting, a mechanism specifically created to combat the target leakage that plagues models trained on correlated data ²⁰⁴ ²⁰⁷. As established, the spatial and temporal proximity of thermal hotspots means that information can "spill over" from training samples to validation samples in a standard cross-validation setup, leading to overly optimistic performance estimates and poor generalization ⁴⁴⁴⁴⁴⁶. Standard boosting algorithms build trees sequentially, where each new tree corrects the errors of the previous ones. However, when calculating the gradient (the direction of improvement) for a new tree, they typically use information from the entire dataset up to that point. This can allow a model to indirectly "see" information about a nearby point's target value, corrupting the learning process. CatBoost's Ordered Boosting addresses this by using a clever permutation trick. During training, it creates multiple ordered versions of the dataset. When building a new tree, it uses one permutation to calculate gradients for all data points, but it calculates the target statistics (like TS) for a given point using only the data points that come *before* it in a *different*, independent permutation ⁴⁴⁸. This mathematical guarantee ensures that the model's prediction for any given sample is based solely on information from other, distinct samples, effectively breaking the chain of leakage. This innovation is not just a theoretical curiosity; it has been shown to significantly improve model generalization, especially in geospatial contexts where spatial dependence is a dominant factor ⁴⁶ ³¹¹. Benchmarks have demonstrated that CatBoost's default settings, which leverage these techniques, often yield superior performance compared to heavily tuned versions of XGBoost and LightGBM, particularly on datasets rich with categorical features ⁶⁶ ³³⁸⁴⁵⁷. The original

authors themselves highlight that CatBoost's combination of these innovations leads to its strong performance across a variety of datasets ²⁰⁸.

The scientific validity of these approaches is well-established in the academic literature. The foundational paper by Prokhorenkova et al. (2018) provides the theoretical underpinnings of both Ordered Target Statistics and Ordered Boosting, positioning them as novel solutions to long-standing problems in gradient boosting ²⁰¹²⁰⁵. Subsequent research has applied and validated CatBoost in diverse domains, reinforcing its effectiveness. For instance, studies on wildfire susceptibility mapping have successfully employed CatBoost to integrate climatic, topographic, and anthropogenic predictors, demonstrating its utility in environmental science ³⁰⁶. Other work has leveraged CatBoost for intrusion detection systems, where its ability to handle categorical network data and prevent leakage is paramount for security applications ⁴¹⁶⁴⁶⁶. Even in fields as varied as crater classification from digital elevation models and predicting casualties in terrorist attacks, CatBoost has proven to be a powerful and reliable tool ²⁶⁴³⁸⁵. This body of evidence confirms that CatBoost is not a niche algorithm but a state-of-the-art GBDT library with a strong theoretical basis and proven track record on complex, real-world structured data. Its design philosophy—prioritizing robustness to leakage and ease of use with categorical data—makes it uniquely suited for the specific demands of the FIRMS hotspot classification task. While other models can be made to work, CatBoost provides these critical capabilities out of the box, aligning perfectly with the project's constraints of speed, reliability, and technical correctness.

NTRO Compliance: Preventing Target Leakage and Delivering Defense-Grade Explanations

For a model to be acceptable to the National Technical Research Organization (NTRO), its performance must be built on a foundation of trust, transparency, and reliability. The two pillars of NTRO compliance in this context are the prevention of spatio-temporal target leakage and the generation of defense-grade, analytically verifiable explanations. CatBoost is uniquely positioned to satisfy both requirements through its core algorithmic design and seamless integration with established explainability frameworks. This dual capability elevates it from a merely accurate model to a trustworthy one, making it the only viable option among state-of-the-art alternatives.

The prevention of target leakage is arguably the most critical requirement for a model intended to generalize to new geographic areas. As thermal anomalies are inherently clustered, a model that simply memorizes the training set's spatial layout will fail when faced with a novel region ^{8 116}. CatBoost's Ordered Boosting mechanism provides a direct and mathematically rigorous solution to this problem ²⁰⁴²⁰⁷. By ensuring that the gradient and target statistics for any given data point are calculated using only a strictly prior subset of the data, CatBoost breaks the causal link that allows leakage to occur ⁴⁴⁶. This is not a heuristic or a simple regularization technique; it is an integral part of the boosting algorithm that guarantees the model learns generalized patterns rather than fitting noise and spatial artifacts. This property is explicitly highlighted in CatBoost's documentation and research papers as a key differentiator from other GBDT implementations ³²⁴⁴⁴⁸. The importance of this feature is underscored by the fact that it directly counters a known pitfall in applying machine learning to geospatial data, where standard cross-validation can produce wildly optimistic and misleading performance metrics ^{46 266}. For an NTRO application, where deploying a model that performs well on old data but fails in the field is unacceptable, this built-in safeguard against leakage is not just a feature—it is a prerequisite. While other models like LightGBM are susceptible to this issue, and XGBoost requires careful management, CatBoost's ordered approach provides a more robust and less error-prone solution, which is crucial for a project managed by a single engineer under tight deadlines ²¹⁵⁴⁴⁸.

The second pillar of NTRO compliance is explainability. The prohibition of black-box models means that every classification must be accompanied by a clear, logical, and mathematically sound justification ¹⁷³²³⁹. CatBoost excels in this domain through its native and highly optimized integration with the `shap.TreeExplainer` module from the SHAP (SHapley Additive exPlanations) library ¹⁸¹²⁴⁵. SHAP is grounded in cooperative game theory and provides a unified measure of feature contribution to a prediction, offering both global insights into model behavior and local explanations for individual instances ¹⁸⁴²⁹¹. For tree-based models like CatBoost, `TreeExplainer` is a fast and exact algorithm that leverages the tree structure to compute SHAP values in polynomial time, avoiding the approximations and computational expense of model-agnostic methods like KernelSHAP ⁴³⁶⁴⁵⁰. This integration produces precise, analytically verifiable outputs. An analyst can examine the SHAP values for a specific hotspot classification and see exactly how each feature contributed to the final decision. For example, the output could definitively state that a classification as "Industrial" was driven by a high `persistence_90d` score, a moderate `frp`, and the absence of nearby water bodies, as captured by OpenStreetMap data. This level of detail is precisely what is needed for forensic analysis and auditability in a defense context ²⁶³³²⁴.

The stability and reliability of SHAP explanations generated for CatBoost have also been validated in various studies. Research indicates that SHAP provides stable feature importance rankings even under conditions of class imbalance, a common issue in classification tasks with unequal class distributions ²⁴⁸³⁹³. This is important for the FIRMS problem, where classes like "Wildfire" might be far less frequent than "Industrial" activity. Furthermore, because CatBoost handles categorical features internally and correctly via Target Statistics, the resulting SHAP values are not contaminated by the artifacts that can arise from flawed external encoding schemes (like the false ordinality imposed by label encoding) ⁸⁵. This ensures that the explanations are faithful to the model's actual learned behavior. While other GBDTs like XGBoost and LightGBM also support TreeExplainer, CatBoost's superior handling of the high-cardinality categorical features (H3 indices, land-use types) means the underlying model is more robust, which in turn leads to more reliable and interpretable SHAP values ¹²⁵²⁸⁸. The combination of Ordered Boosting for trustworthy predictions and native SHAP integration for transparent explanations makes CatBoost the only model that fully satisfies the stringent NTRO requirements. It delivers not just accuracy, but accountable and auditable accuracy.

Comparative Analysis: A Brutal Evaluation of Alternative Models

While CatBoost emerges as the leading candidate, a rigorous adversarial evaluation requires subjecting it to the same scrutiny applied to its powerful competitors. Models like XGBoost and LightGBM are industry standards and often win Kaggle competitions, making them natural contenders ^{63 65}. Similarly, deep learning architectures like TabNet and spatially-aware models like Graph Neural Networks (GNNs) represent the cutting edge of AI. However, a closer examination reveals fundamental flaws in each of these alternatives when measured against the specific constraints of the FIRMS classification task: high-cardinality spatial data, CPU-only training, and NTRO-level explainability.

Model Family	Primary Strength(s)	Critical Failure(s) for FIRMS Task
XGBoost	Robust baseline performance, flexibility, wide community support ^{84 438} .	Vulnerable to target leakage in spatial data ²¹⁵ ; requires problematic external encoding (OHE/LB) for H3 indices, risking memory overflow or false assumptions ^{85 298} .
LightGBM	High computational speed and efficiency ¹⁶²¹⁶⁵ .	Prone to overfitting and target leakage due to leaf-wise tree growth; lacks CatBoost's rigorously designed leakage prevention mechanisms ²¹⁸⁴⁴⁸ .
Deep Learning (TabNet)	State-of-the-art accuracy on some tabular benchmarks ⁴⁷ .	Requires GPU acceleration for feasible training times ⁴⁷⁸ ; inherently "black-box," making it difficult to generate stable, interpretable SHAP explanations required by NTRO ^{21 156} .
Graph Neural Networks (GNNs)	Designed for relational/spatial data, capturing neighborhood effects ⁸³ .	Over-engineering for this problem; requires constructing complex adjacency graphs from H3 grids, writing custom code, and faces convergence issues on CPU ^{23 111} .

XGBoost, often considered the default for GBDT tasks, is a powerful but flawed choice here. Its primary weakness is its vulnerability to target leakage, a known issue in geospatial modeling ²¹⁵. Without specific safeguards, it can overfit to the spatial structure of the training data, leading to poor generalization. More critically, XGBoost has no native way to handle the H3 hexagon index. The two standard encoding methods are inadequate. One-Hot Encoding (OHE) would be disastrous, creating a feature matrix far too large for any CPU environment to handle, leading to immediate OOM errors ²⁹⁸. Label Encoding, the alternative, is equally problematic as it imposes a false ordinal relationship on the categorical values. The model would incorrectly assume that there is a meaningful mathematical order between H3 cell IDs (e.g., that cell ID 8828308 is "greater than" cell ID 8828301), which is nonsensical and would corrupt the model's ability to learn spatial relationships correctly. CatBoost's internal Target Statistics encoding avoids both of these pitfalls entirely, making it a far more robust and appropriate choice for this data structure ⁸⁵.

LightGBM is renowned for its speed and efficiency, often outperforming XGBoost in terms of training time ¹⁶²¹⁶⁵. However, its very strengths can become liabilities in this context. Its leaf-wise tree growth strategy, while efficient, can lead to overfitting on noisy or correlated data ¹⁶⁰. More importantly, it shares XGBoost's reliance on external categorical encodings, facing the same OHE and label encoding dilemmas. While LightGBM has some built-in categorical support, CatBoost's Ordered Boosting mechanism is more explicitly and rigorously designed to prevent the specific type of spatio-temporal leakage that is the central concern of this task ²¹⁸⁴⁴⁸. Given the 3-week timeline, implementing and validating the necessary custom fixes for LightGBM to make it safe for this spatial data would be a significant risk and likely consume more time than starting with CatBoost, which is designed to be safe by default.

Deep Learning Models, such as those based on Transformers or TabNet, are often marketed as the future of tabular data analysis ^{18 64}. However, they are fundamentally unsuited for this project's constraints. First, and most practically, they do not train on CPUs in a reasonable timeframe; GPU acceleration is almost always required, which is forbidden by the Kaggle environment constraint ⁴⁷⁸. Second, and more philosophically, they fail the NTRO's explainability requirement. Their internal mechanisms, such as attention masks, do not lend themselves to clean, stable, and interpretable SHAP values in the way that tree-based models do ^{21 156}. They function as black boxes, producing predictions without providing the analytically verifiable, forensic-level explanations that are non-negotiable for a defense application ²³⁹. Attempting to force a deep learning model into this role would be a strategic failure, regardless of its potential accuracy.

Finally, **Graph Neural Networks (GNNs)** and **Spatial CNNs** represent another category of over-engineering. The problem can be solved by aggregating data to H3 cells and then treating each cell as an independent tabular datapoint ¹¹¹. There is no need to build a complex graph where each node is an individual hotspot and edges represent spatial proximity. Such an approach would require manually constructing an H3 adjacency graph, which is computationally intensive, and writing custom PyTorch Geometric dataloaders, which is a significant software engineering task ^{23 83}. This path is fraught with complexity and risk, especially on a CPU, and would likely fail to converge or produce meaningful results within the 3-week deadline. The problem is better framed as a classification task on aggregated spatial features, a domain where GBDTs excel. Therefore, GNNs are not just unnecessary; they are an inappropriate and inefficient solution to the problem as posed.

In conclusion, the brutal evaluation of alternatives reveals why CatBoost is the only viable choice. XGBoost and LightGBM are compromised by their handling of categorical data and leakage. Deep learning models are infeasible due to hardware and explainability constraints. GNNs are a case of solving the wrong problem with excessive complexity. CatBoost, by contrast, offers a balanced solution that is powerful, robust, feasible, and compliant with all specified requirements.

Actionable Implementation Plan for CatBoost

To successfully implement the CatBoost model for the NASA FIRMS hotspot classification task, a clear, step-by-step plan is required. This plan addresses data preparation, model configuration, training, evaluation, and explanation generation, all tailored to the specific

challenges of high-cardinality spatial data and the NTRO's stringent requirements. Adhering to this plan will maximize the chances of building a robust, trustworthy, and operationally feasible model within the 3-week timeline.

Step 1: Data Preparation and Feature Engineering The most critical aspect of preparation is to avoid any premature or incorrect encoding of the high-cardinality categorical features. The H3 cell identifier, OpenStreetMap-derived land-use tags, and satellite instrument names (SNPP, NOAA-20) must be passed to the CatBoost model as-is, typically as string data types. Do not apply One-Hot Encoding or any form of label encoding to these columns. CatBoost's strength lies in its ability to handle them natively. The focus should instead be on engineering the other features. The `type` field, if available, should be carefully examined; it appears to be an inferred label rather than a perfectly reliable one, so it may be better used as a feature (`is_static_source`) rather than a ground-truth label to avoid introducing potential bias ²⁰². Features derived from date/time, such as `persistence_90d`, must be normalized correctly, accounting for any known gaps in the data to avoid spuriously low persistence scores near those gaps ⁴⁶. For instance, the 11 missing days found in the EDA phase must be handled by normalizing the persistence count by the number of observed days in the 90-day window, not by a fixed divisor of 90 ¹.

Step 2: Model Initialization and Configuration Initialize the CatBoost classifier with parameters that are well-suited for this task. The key parameter is `cat_features`, a list containing the names of the columns that contain categorical data. This tells CatBoost which columns to process with its specialized Target Statistics and Ordered Boosting algorithms. Other important parameters include:

- `eval_metric='MultiClass'`: Since the task involves six classes, this metric will be used to monitor performance during training and validation.
- `verbose=100`: Set to a positive integer to print training progress periodically, which is useful for monitoring convergence.
- `random_seed=42`: Set a fixed random seed for reproducibility.
- `auto_class_weights='Balanced'`: This is a crucial parameter for handling the likely class imbalance across the six classes (e.g., Industrial vs. Mining vs. Wildfire). This setting automatically adjusts the weights of the classes during training to give more importance to minority classes, improving the model's ability to detect them ³⁸⁹³⁹⁰. This is a much simpler and often effective approach than manually tuning `scale_pos_weight`, which is more complex for multi-class problems ³⁹².

Step 3: Training and Validation Strategy Standard K-Fold Cross-Validation is inappropriate for this dataset due to the strong spatial autocorrelation. Using it would

likely result in overly optimistic performance metrics because the same geographic regions could appear in both training and testing folds ⁴⁶. The correct approach is **Spatial Cross-Validation (SCV)**. SCV is a resampling method designed to mimic real-world deployment by ensuring that training and testing sets are separated by a minimum distance. This prevents information from "spilling over" from the training area to the prediction area, providing a more realistic and pessimistic estimate of the model's out-of-sample performance ^{46 116}. Implementing SCV can be done using libraries like `geopy` to calculate distances between fold centroids or by using dedicated packages like `scikit-learn-extra`. This step is non-negotiable for producing a model that the NTRO can trust.

Step 4: Hyperparameter Tuning While CatBoost is known for performing well out of the box, some hyperparameter tuning will likely be necessary to reach peak performance ⁴⁵⁷. Use a Bayesian optimization library like Optuna or a randomized search within the SCV framework to find optimal parameters. Focus tuning on key parameters such as `learning_rate`, `depth` (tree depth), `l2_leaf_reg` (L2 regularization), and `bagging_temperature` (for stochastic gradient boosting). Avoid tuning parameters related to categorical encoding, as CatBoost's internal methods are already optimal for this purpose.

Step 5: Explanation Generation with SHAP After training the final model, generate SHAP values to fulfill the NTRO's explainability requirement. CatBoost provides a convenient method for this:

```
# Assuming 'model' is your trained CatBoostClassifier and 'test_pool' is a
shap_values = model.get_feature_importance(type='ShapValues', data=test_po
```

This will return a matrix of SHAP values with dimensions `(n_test_samples, n_features + 1)` for a multi-class problem. The extra column represents the impact of the model's base output (the expected value). These values can then be visualized using various SHAP plots. Summary plots can show the overall most important features driving the classifications, while force plots can provide a detailed, per-instance explanation for a specific hotspot, showing exactly how each feature pushed the prediction towards a particular class ¹⁷⁷¹⁷⁸. These visualizations are the "defense-grade" output the NTRO requires and should be integrated into the final reporting dashboard.

By following this structured plan, the single ML engineer can systematically build, validate, and explain a state-of-the-art model that is tailored to the unique challenges of the FIRMS dataset and meets all the operational and ethical constraints of the project.

Final Verdict: Justification and Strategic Recommendation

After a comprehensive and adversarial evaluation of the problem constraints, data characteristics, and candidate machine learning models, the final verdict is unequivocal. For the task of classifying NASA FIRMS thermal hotspots into six distinct classes under the specified operational constraints, **CatBoost is the single best model**. Its selection is not based on generic popularity or minor performance gains, but on its unique and mathematically-grounded solutions to the three core challenges of this problem: the high-cardinality nature of spatial data, the pervasive threat of spatio-temporal target leakage, and the non-negotiable requirement for defense-grade explainability.

The justification for this recommendation rests on three pillars. First, CatBoost's native handling of categorical features via **Target Statistics** provides an elegant and memory-efficient solution to the challenge of millions of H3 hexagon indices. Unlike XGBoost or LightGBM, which force a choice between memory-prohibitive One-Hot Encoding and logic-corrupting Label Encoding, CatBoost automates a statistically sound encoding process that preserves the integrity of the data while keeping the feature space manageable ^{85 215}. This directly addresses the CPU-only training constraint and saves invaluable development time. Second, CatBoost's **Ordered Boosting** mechanism offers a rigorous, algorithmic defense against spatio-temporal target leakage. By ensuring that a model's prediction for any given point is informed only by strictly prior data in a permuted sequence, it prevents the model from memorizing the training map and forces it to learn generalizable, physically meaningful patterns ²⁰⁷⁴⁴⁶. This is essential for building a model that will perform reliably in new geographic deployments, a critical requirement for any NTRO-backed system. Third, CatBoost's seamless integration with the **SHAP framework** provides the exact kind of analytically verifiable, forensic-level explanations demanded by intelligence agencies. The ability to generate precise, game-theory-backed attribution of each feature's contribution to a final classification transforms the model from a black box into a transparent and auditable decision-support tool ¹⁸¹²⁹¹.

All viable alternatives were subjected to a brutal rejection process. XGBoost and LightGBM, while powerful, are fundamentally compromised by their handling of categorical data and their susceptibility to leakage without extensive, risky custom modifications. Deep learning models like TabNet are disqualified by their GPU dependency and inherent lack of transparency. More complex architectures like Graph Neural Networks are a case of severe over-engineering for a problem that is more naturally framed as a classification task on aggregated tabular data. No other model

offers the same combination of out-of-the-box capabilities that are perfectly aligned with the project's specific needs.

Therefore, the strategic recommendation is to proceed with CatBoost. The implementation plan outlined—a process of careful data preparation, targeted model configuration with balanced class weights, rigorous spatial cross-validation, and systematic explanation generation—is the definitive path forward. By adopting CatBoost, the project can deliver a solution that is not only accurate and efficient but, most importantly, trustworthy, transparent, and compliant with the highest standards of defense technology evaluation.

Reference

1. Challenges in data-driven geospatial modeling for ... - PMC <https://pmc.ncbi.nlm.nih.gov/articles/PMC11659326/>
2. Incorporating Spatial Autocorrelation in Machine Learning ... <https://www.mdpi.com/2220-9964/11/4/242>
3. Geospatial SHAP interpretability for urban road collapse ... <https://www.tandfonline.com/doi/full/10.1080/19475705.2025.2491473>
4. The problematic case of data leakage: A case for leave- ... <https://www.sciencedirect.com/science/article/pii/S0016706125000618>
5. Data-Driven Geospatial Modeling and Forecasting of ... <https://www.medrxiv.org/content/10.64898/2026.08.16.26360539v1.full>
6. Leveraging the TabPFN Algorithm for High - Resolution ... <https://agupubs.onlinelibrary.wiley.com/doi/full/10.1029/2025WR042780>
7. Identifying associated variables using Shapley values | BMC ... <https://link.springer.com/article/10.1186/s12889-026-26571-6>
8. 13 Geospatial Modeling https://warin.ca/geospatial/13-geospatial_models.html
9. interpretable HGB-SHAP machine learning model <https://www.frontiersin.org/journals/environmental-science/articles/10.3389/fenvs.2025.1640840/full>
10. Making geo joins faster with H3 indexes <https://news.ycombinator.com/item?id>
11. Unlock the Potential of Spatial Indexes: 10 Powerful Uses ... <https://carto.com/blog/h3-spatial-indexes-10-use-cases/>

12. Unlocking Spatial Intelligence with H3: A Practitioner's ... <https://medium.com/codex/unlocking-spatial-intelligence-with-h3-a-practitioners-guide-to-hexagonal-hierarchical-spatial-45be909008ae>
13. r/gis - How H3 Hexagons Turn Geography into Drive-Time ... https://www.reddit.com/r/gis/comments/1sumk4w/how_h3_hexagons_turn_geography_into_drivetime/
14. Simplify Spatial Indexing with the Power of H3 <https://www.snowflake.com/en/blog/getting-started-with-h3-hexagonal-grid/>
15. Performance characteristics of index init · Issue #168 · uber ... <https://github.com/uber/h3/issues/168>
16. Too slow or out of memory problems in Machine Learning/ ... <https://www.kaggle.com/general/4024>
17. Amazon Redshift, CARTO, and H3 | AWS Big Data Blog <https://aws.amazon.com/blogs/big-data/breaking-barriers-in-geospatial-amazon-redshift-carto-and-h3/>
18. Tabular Data: Is Deep Learning all you need? <https://arxiv.org/html/2402.03970v3>
19. Spatio-Temporal Missing Data Imputation: A Systematic ... <https://dl.acm.org/doi/full/10.1145/3797903>
20. Machine Learning's Triumph in Water Quality Forecasting <https://arxiv.org/html/2309.16951v2>
21. Spatio-Temporal Predictive Modeling Techniques for ... <https://dl.acm.org/doi/10.1145/3696661>
22. Representation Learning for Tabular Data <https://arxiv.org/html/2504.16109v1>
23. Deep Learning for Spatio-Temporal Data Mining: A Survey <https://arxiv.org/pdf/1906.04928>
24. Spatio-Temporal Foundation Models: Vision, Challenges, ... <https://arxiv.org/html/2501.09045v2>
25. Deep Learning within Tabular Data: Foundations ... <https://arxiv.org/html/2501.03540v1>
26. Interpretable Spatio-Temporal Features Extraction based ... <https://arxiv.org/html/2506.00858v1>
27. Large Models for Time Series and Spatio-Temporal Data <https://dl.acm.org/doi/full/10.1145/3821637>
28. (PDF) Synergetic Use of Random Forest and H3 Grid Cell ... https://www.researchgate.net/publication/391616640_Synergetic_Use_of_Random_Forest_and_H3_Grid_Cell_for_the_Flood_Monitoring_Using_Sentinel-1_Data

29. Urban low-altitude airspace resources: Integrating risk and ... <https://www.sciencedirect.com/science/article/pii/S0264275126004221>
30. a foundation model approach for disaster detection from social ... <https://link.springer.com/article/10.1007/s11069-026-08191-w>
31. Talks <https://2026.europe.foss4g.org/schedule/talks/>
32. Understanding and Predicting Global Starlink Performance <https://www.nitindermohan.com/documents/2026/pubs/sigmetrics26-horizon.pdf>
33. Unsupervised & Semi-Supervised Learning https://www.climatechange.ai/subject_areas/unsupervised_semi_supervised_learning
34. AI for detecting and localizing concurrent abrupt ... <https://pubmed.ncbi.nlm.nih.gov/41707605/>
35. ISRS 2026 Technical Program https://www.rssj.or.jp/isrs2026/program_April_14.html
36. ICRA 2026 Program | Wednesday June 3, 2026 https://ras.papercept.net/conferences/conferences/ICRA26/program/ICRA26_ContentListWeb_4.html
37. FIRMS | NASA Earthdata <https://www.earthdata.nasa.gov/data/tools/firms>
38. Satellite (VIIRS) Thermal Hotspots and Fire Activity <https://data-nifc.opendata.arcgis.com/datasets/esri2::satellite-viirs-thermal-hotspots-and-fire-activity/about>
39. Archival NRT FIRMS Global VIIRS and MODIS vector data https://gee-community-catalog.org/projects/firms_vector/
40. Assessing the performance of MODIS and VIIRS active fire ... <https://iforest.sisef.org/contents/?id>
41. NASA FIRMS - Active Fire https://pyronear.org/pyro-risks/overview/datasets/NASA-FIRMS_ACTIVE-FIRE_VIIRS.html
42. FIRMS: Fire Information for Resource Management System <https://developers.google.com/earth-engine/datasets/catalog/FIRMS>
43. NASA FIRMS Active Fire Data (MODIS/VIIRS) <https://www.kaggle.com/datasets/vijayveersingh/nasa-firms-active-fire-dataset-modisviirs>
44. Geospatial system for wildfire monitoring and prediction ... <https://www.sciencedirect.com/science/article/pii/S2666592125001167>
45. Geographically weighted machine learning for predicting ... <https://www.sciencedirect.com/science/article/pii/S2214581825006962>
46. Accounting for Spatial Autocorrelation in Algorithm-Driven ... <https://link.springer.com/article/10.1007/s11146-022-09915-y>

47. Tabular Data Classification and Regression : XGBoost or ... https://www.researchgate.net/publication/387099599_Tabular_Data_Classification_and_Regression_XGBoost_or_Deep_Learning_with_Retrieval-Augmented_Generation
48. Landslide susceptibility modeling based on SHAP ... <https://www.frontiersin.org/journals/earth-science/articles/10.3389/feart.2025.1731872/full>
49. An Efficient Transformer Model for Geo-Spatial Tabular Data <https://ojs.aaai.org/index.php/AAAI/article/view/33259/35414>
50. 13 Beyond Regression and Classification - mlr3 book <https://mlr3book.mlr-org.com/survival.html>
51. SHAP-Based Spatial Decomposition of Crop Yield ... <https://www.ijert.org/shap-based-spatial-decomposition-of-crop-yield-prediction-errors-across-uttar-pradesh-s-agro-climatic-zones-a-district-level-explainable-machine-learning-framework-ijertv15is060737>
52. Do foundation models work for geospatial tabular data? An ... <https://www.tandfonline.com/doi/full/10.1080/13658816.2026.2691066>
53. present VIIRS NOAA-20 Temporal Coverage: 1 January 2020 https://firms.modaps.eosdis.nasa.gov/content/descriptions/FIRMS_VIIRS_Firehotspots.html
54. FIRMS FAQ <https://www.earthdata.nasa.gov/data/tools/firms/faq>
55. VIIRS (Suomi NPP, NOAA-20 and ... https://firms.modaps.eosdis.nasa.gov/descriptions/FIRMS_VIIRS_Firehotspots.html
56. Active Fire Data Attributes for MODIS and VIIRS <https://www.earthdata.nasa.gov/data/tools/firms/active-fire-data-attributes-modis-viirs>
57. Archive Download - NASA | LANCE <https://firms.modaps.eosdis.nasa.gov/download/>
58. Near Real-Time Global Fire Monitoring Using Data from ... <https://ntrs.nasa.gov/citations/20190032007>
59. Satellite (VIIRS) Thermal Hotspots and Fire Activity <https://atlas.eia.gov/datasets/esri2::satellite-viirs-thermal-hotspots-and-fire-activity/about>
60. Satellite (VIIRS) Thermal Hotspots and Fire Activity - Datasets <https://wifire-data.sdsc.edu/dataset/satellite-viirs-thermal-hotspots-and-fire-activity>
61. 0.97952 - LightGBM + XGBoost + CatBoost Ensemble <https://www.kaggle.com/competitions/tabular-playground-series-may-2022/writeups/james-mcguigan-541th-place-0-97952-lightgbm-xgboos>
62. Benchmarking TabPFN V2 against XGBoost and CatBoost ... <https://medium.com/@humblebeeai-team/benchmarking-tabpfv2-against-xgboost-and-catboost-on-kaggle-datasets-7e199dfd9f77>

63. Tabular Machine Learning 2026: TabArena Leaderboard <https://www.codesota.com/tasks/tabular-ml>
64. Foundation Models for Tabular Data – The Future of ... <https://www.kaggle.com/general/566353>
65. When Do Neural Nets Outperform Boosted Trees on ... <https://arxiv.org/pdf/2305.02997>
66. XGBoost vs LightGBM vs CatBoost vs AdaBoost <https://www.kaggle.com/code/faressayah/xgboost-vs-lightgbm-vs-catboost-vs-adaboost>
67. A Comprehensive Benchmark of Machine and Deep ... <https://arxiv.org/html/2408.14817v1>
68. How TabPFN Leverages In-Context Learning to Achieve ... <https://www.marktechpost.com/2026/04/19/how-tabpfn-leverages-in-context-learning-to-achieve-superior-accuracy-on-tabular-datasets-compared-to-random-forest-and-catboost/>
69. Our new classification algorithm outperforms CatBoost ... <https://www.kaggle.com/general/502240>
70. Forest Fires - Burnt Area Prediction <https://www.kaggle.com/c/forestfiresarea/data>
71. NASA FIRMS Active Fire Data (MODIS/VIIIRS) <https://www.kaggle.com/datasets/vijayveersingh/nasa-firms-active-fire-dataset-modisviirs/data>
72. Machine Learning for Data Science VII <https://www.kaggle.com/c/machine-learning-for-data-science-vii/data>
73. Find Open Datasets for AI and Research <https://www.kaggle.com/datasets?search>
74. FLAME 3 -Nadir Thermal Fire Image Dataset <https://www.kaggle.com/datasets/brycehopkins/flame-3-nadir-thermal-plot-subset>
75. Wildfire Prediction Dataset (Satellite Images) <https://www.kaggle.com/datasets/abdelghaniaaba/wildfire-prediction-dataset>
76. The wildfire dataset <https://www.kaggle.com/datasets/elmadafri/the-wildfire-dataset/data>
77. TS-SatFire <https://www.kaggle.com/datasets/z789456sx/ts-satfire>
78. FLIR Thermal Images Dataset [reduced] <https://www.kaggle.com/datasets/albertofv/flir-thermal-images-dataset-reduced>
79. Towards eXplainable AI for Mobility Data Science <https://arxiv.org/html/2307.08461v3>
80. A Systematic Review of Spatio-Temporal Statistical Models <https://arxiv.org/html/2511.00422v1>

81. A Systematic Review on Explainable AI in Remote Sensing <https://arxiv.org/html/2402.13791v1>
82. Interpretable Machine Learning for Weather and Climate ... <https://arxiv.org/html/2403.18864v1>
83. A Systematic Literature Review of Spatio-Temporal Graph ... <https://arxiv.org/html/2410.22377v2>
84. Which is best catboost, xgboost, lightgbm which have ... <https://www.kaggle.com/questions-and-answers/512218>
85. XGBoost, LightGBM, and Other Kaggle Competition Favorites <https://medium.com/analytics-vidhya/xgboost-lightgbm-and-other-kaggle-competition-favorites-6212e8b0e835>
86. CatBoost, LightGBM, XGBoost: CPU vs GPU <https://www.kaggle.com/code/marychin/catboost-lightgbm-xgboost-cpu-vs-gpu>
87. TS-SatFire: A Multi-Task Satellite Image Time-Series Dataset ... <https://pmc.ncbi.nlm.nih.gov/articles/PMC12630597/>
88. Satellite-Based Wildfire Monitoring Across the Fire ... <https://www.kjrs.org/journal/view.html?uid>
89. Satellite Observations and Tools for Fire Risk, Detection ... <https://www.earthdata.nasa.gov/learn/trainings/satellite-observations-tools-fire-risk-detection-analysis>
90. Improving Wildfire Tracking with Satellite Detection - ESSIC <https://essic.umd.edu/improving-wildfire-tracking-with-satellite-detection/>
91. Assessment and monitoring of post-fire recovery using ... <https://ui.adsabs.harvard.edu/abs/2025EGUGA..27..183Z/abstract>
92. Wildfire Detection And Forest Fire Monitoring With Satellite ... <https://eos.com/wildfire-detection-and-monitoring/>
93. Machine Learning Approach to Improve Monitoring Active ... <https://ui.adsabs.harvard.edu/abs/2024AGUFMB33I...07L/abstract>
94. An Efficient Wildfire Detection System for AI-Embedded ... <https://www.mdpi.com/2571-6255/6/4/169>
95. Revolutionizing Wildfire Monitoring with ICEYE's SAR ... <https://www.iceye.com/blog/wildfire-monitoring-with-sar>
96. Software Engineering: Emerging Trends and Practices in ... <https://www.scribd.com/document/991865690/Software-Engineering-Emerging-Trends-and-Practices-in-System-Development-Proceedings-of-14th-Computer-Science-on-line-Radek-Silhavy-Petr-Silhavy>

97. Satellite (VIIRS) Thermal Hotspots and Fire Activity <https://www.arcgis.com/home/item.html?id>
98. Active Fire Monitoring with NASA FIRMS: Real-Time ... <https://offnadir-delta.com/blog/nasa-firms-active-fire-monitoring>
99. Editorial: Understanding NASA's Hotspot Maps ... <https://www.facebook.com/MinistryOfEnvironmentCambodia/posts/editorial-understanding-nasas-hotspot-maps-and-the-risks-of-misinterpretationfeb/1237253215262828/>
100. Satellite (VIIRS) Thermal Hotspots and Fire Activity <https://climate.esri.ca/search?tags>
101. Assessing the performance of MODIS and VIIRS active fire ... <https://iforest.sisef.org/pdf/?id>
102. Rapid Wildfire Hotspot Detection Using Self-Supervised Learning on ... <https://arxiv.org/html/2405.20093v1>
103. AI Geospatial Wildfire Risk Prediction | Kaggle <https://www.kaggle.com/code/thjaquenoud/ai-geospatial-wildfire-risk-prediction>
104. TS-SatFire: A Multi-Task Satellite Image Time-Series Dataset ... - Nature <https://www.nature.com/articles/s41597-025-06271-3>
105. A dataset of multi-modal time series for wildfire spread prediction <https://openreview.net/forum?id=RgdGkPRQ03>
106. Geospatial Data Analytics (Competition I) - Kaggle <https://www.kaggle.com/competitions/geospatial-data-analytics>
107. FlameGuard: an AI-driven smart map for early detection and ... - PMC - NIH <https://pmc.ncbi.nlm.nih.gov/articles/PMC12663161/>
108. a comprehensive benchmark dataset for sub-national crop yield forecasting <https://essd.copernicus.org/articles/18/3997/2026/>
109. A Machine Learning Dataset to Predict Wildfire Spreading From Remote ... https://www.researchgate.net/publication/362272369_Next_Day_Wildfire_Spread_A_Machine_Learning_Dataset_to_Predict_Wildfire_Spreading_From_Remote-Sensing_Data
110. Geospatial Feature Engineering and Visualization - Kaggle <https://www.kaggle.com/code/camnugent/geospatial-feature-engineering-and-visualization>
111. H3-MOSAIC: multimodal generative AI for semantic place detection ... - PMC <https://pmc.ncbi.nlm.nih.gov/articles/PMC12640565/>
112. [PDF] A Planet Scale Spatial-Temporal Knowledge Graph Based On ... <https://ceur-ws.org/Vol-3743/paper6.pdf>
113. Beyond Square Pixels: H3 Spatial Indexing for Global Raster Data <https://www.youtube.com/watch?v=IUrh9NtSe6s>

114. Built-in H3 for Geospatial Analytics | Databricks Blog <https://www.databricks.com/blog/announcing-built-h3-expressions-geospatial-processing-and-analytics>
115. Introduction | H3 <https://h3geo.org/docs/>
116. From Machine Learning to Large-Scale EO Products - arXiv <https://arxiv.org/html/2607.24532v3>
117. (PDF) An Enhanced Earthquake Risk Analysis using H3 Spatial Indexing https://www.researchgate.net/publication/374458826_An_Enhanced_Earthquake_Risk_Analysis_using_H3_Spatial_Indexing
118. Understanding spatial indexes: H3 explained <https://felt.com/blog/h3-spatial-index-hexagons>
119. H3 Indexing for Massive Spatial Datasets <https://community.databricks.com/t5/technical-blog/customer-blog-scaling-geospatial-analytics-efficient/ba-p/131234>
120. Fast Geospatial Indexing with H3 <https://medium.com/data-science/fast-geospatial-indexing-with-h3-90e862482585>
121. A Survey on Spatio-Temporal Prediction <https://dl.acm.org/doi/10.1145/3766546>
122. Foundation Models for Spatio-Temporal Data Science <https://arxiv.org/html/2503.13502v1>
123. Real-Time Explanations for Tabular Foundation Models <https://arxiv.org/html/2603.29946v1>
124. A Closer Look at Deep Learning Methods on Tabular ... <https://arxiv.org/html/2407.00956v4>
125. Tabular examples — SHAP latest documentation https://shap.readthedocs.io/en/latest/tabular_examples.html
126. Benchmark XGBoost explanations - the SHAP documentation https://shap.readthedocs.io/en/latest/example_notebooks/benchmarks/tabular/Benchmark%20XGBoost%20explanations.html
127. A Closer Look at Deep Learning Methods on Tabular ... <https://arxiv.org/html/2407.00956v2>
128. Benchmarks — SHAP latest documentation <https://shap.readthedocs.io/en/latest/benchmarks.html>
129. (PDF) A Comprehensive Benchmark of Machine and Deep ... https://www.researchgate.net/publication/383461054_A_Comprehensive_Benchmark_of_Machine_and_Deep_Learning_Across_Diverse_Tabular_Datasets
130. Deep Learning Benchmark for Tabular Data | PDF <https://www.scribd.com/document/781254045/A-Closer-Look-at-Deep-Learning-on-Tabular-Data>

131. what's wrong with your synthetic tabular data? using ... <https://arxiv.org/pdf/2504.20687>
132. NASA Active Fire Product - VIIRS Land Team <https://viirsland.gsfc.nasa.gov/Products/NASA/FireESDR.html>
133. VIIRS Global Burned Area Product Added to FIRMS <https://www.earthdata.nasa.gov/news/blog/viirs-global-burned-area-product-added-firms>
134. FIRMS Adds Ultra Real-Time Data from MODIS and VIIRS <https://www.earthdata.nasa.gov/news/feature-articles/firms-adds-ultra-real-time-data-from-modis-viirs>
135. The NOAA VIIRS Active Fire Product https://www.star.nesdis.noaa.gov/star/documents/seminardocs/2022/20220420_VAWS_Csiszar_VIIRS_fire.pdf
136. Characteristics of VIIRS, MODIS, and OLI Sensors and their ... <https://www.earthdata.nasa.gov/news/blog/characteristics-viirs-modis-oli-sensors-effects-spatial-extent-daily-active-fire-data>
137. VIIRS - Dark Target <https://darktarget.gsfc.nasa.gov/satellite/viirs>
138. What caveats should be considered when using active fire ... <https://forum.earthdata.nasa.gov/viewtopic.php?t>
139. catboost/benchmarks: Comparison tools <https://github.com/catboost/benchmarks>
140. CatBoost vs. XGBoost: Why CatBoost is Stealing the Crown in ... <https://valeman.medium.com/catboost-vs-xgboost-why-catboost-is-stealing-the-crown-in-tabular-data-2d1c2f96afbc>
141. Realistic Evaluation of TabPFN v2 in Open Environments <https://arxiv.org/html/2505.16226v1>
142. Benchmarking state-of-the-art gradient boosting algorithms ... <https://arxiv.org/pdf/2305.17094>
143. 1 TabPFN-2.5 performance on the standard TabArena-lite ... <https://arxiv.org/html/2511.08667v1>
144. Transformers Boost the Performance of Decision Trees on ... <https://arxiv.org/html/2502.02672v1>
145. TabPFN vs XGBoost: Performance degradation with large ... <https://github.com/PriorLabs/TabPFN/issues/257>
146. Where CatBoost beats XGBoost and LightGBM <https://valeman.substack.com/p/where-catboost-beats-xgboost-and>
147. Comparative analysis of CatBoost, LightGBM, XGBoost, RF ... https://www.researchgate.net/publication/391493901_Comparative_analysis_of_CatBoost_LightGBM_XGBoost_RF_and_DT_met

hods_optimised_with_PSO_to_estimate_the_number_of_k-
barriers_for_intrusion_detection_in_wireless_sensor_networks

148. Deep Learning for Time Series Classification and Extrinsic ... <https://dl.acm.org/doi/10.1145/3649448>
149. Computer Science <https://arxiv.org/list/cs/new>
150. Context-Aware Multimodal Representation Learning for ... <https://arxiv.org/html/2511.11706v1>
151. Earth Embeddings <https://arxiv.org/pdf/2608.03410>
152. Bayesian Inference of Nonlinear Malaria Dynamics in ... <https://arxiv.org/html/2606.00783v1>
153. Context-Aware Multimodal Representation Learning for ... <https://ieeexplore.ieee.org/iel8/36/11333908/11585830.pdf>
154. Computer Vision and Pattern Recognition <https://arxiv.org/list/cs.CV/new>
155. From Deep Learning to LLMs: A survey of AI in Quantitative ... <https://arxiv.org/html/2503.21422v1>
156. SHAP-DNN Approach Advances Remote Sensing ... <https://ieeexplore.ieee.org/iel8/4609443/10766875/10884714.pdf>
157. Video Understanding by Design: How Datasets Shape ... <https://arxiv.org/html/2509.09151v2>
158. A Comparative Study of XGBoost, LightGBM, and CatBoost Models for ... https://www.researchgate.net/publication/397440582_A_Comparative_Study_of_XGBoost_LightGBM_and_CatBoost_Models_for_Customer_Churn_Prediction_in_the_Banking_Industry
159. 771: Gradient Boosting: XGBoost, LightGBM and CatBoost - YouTube <https://www.youtube.com/watch?v=tqR8jGm6Zhw>
160. CatBoost vs. LightGBM vs. XGBoost - Towards Data Science <https://towardsdatascience.com/catboost-vs-lightgbm-vs-xgboost-c80f40662924/>
161. Why Do Tree-Based Models (LightGBM, XGBoost, CatBoost ... https://www.reddit.com/r/learnmachinelearning/comments/1kmdils/why_do_treebased_models_lightgbm_xgboost_catboost/
162. XGBoost vs LightGBM vs CatBoost vs AdaBoost - Kaggle <https://www.kaggle.com/code/faressayah/xgboost-vs-lightgbm-vs-catboost-vs-adaboost/notebook>
163. Ensemble Methods - Boosting (XGBoost, LightGBM, CatBoost) <https://learn.modernagecoders.com/resources/ai-and-machine-learning/ensemble-methods-and-boosting>
164. Ensemble Machine Learning of Gradient Boosting (XGBoost ... - PMC <https://pmc.ncbi.nlm.nih.gov/articles/PMC10611362/>

165. 13 LightGBM and CatBoost – Data Science III with python (Class notes) https://nustat.github.io/STAT303-3-class-notes/Lec11_More%20boosting%20models.html
166. XGBoost, LightGBM or CatBoost - which boosting algorithm should I ... <https://www.riskified.com/resources/article/boosting-comparison/>
167. A Comparative Study of SHAP and LIME in Intrusion Detection Models <https://www.mdpi.com/2076-3417/15/13/7329>
168. Explainable AI for Digital Forensics: Ensuring Transparency in Legal ... <https://www.forensicscijournal.com/journals/jfsr/jfsr-aid1089.php>
169. Explainable AI for forensic speech authentication within cognitive and ... <https://www.frontiersin.org/journals/neuroscience/articles/10.3389/fnins.2025.1692122/full>
170. Explainable AI for forensic speech authentication within cognitive and ... <https://pmc.ncbi.nlm.nih.gov/articles/PMC12626942/>
171. (PDF) Developing an Explainable AI System for Digital Forensics https://www.researchgate.net/publication/394666181_Developing_an_Explainable_AI_System_for_Digital_Forensics_Enhancing_Trust_and_Transparency_in_Flagging_Events_for_Legal_Evidence
172. AI Transparency and Explainability: XAI Techniques and Tools <https://agility-at-scale.com/ai/governance/transparency-and-explainability/>
173. NIST Publishes Proposed Principles for “Explainable” AI Systems <https://www.hunton.com/privacy-and-cybersecurity-law-blog/nist-publishes-proposed-principles-for-explainable-ai-systems>
174. XAI-CF – Examining the Role of Explainable Artificial Intelligence in ... <https://arxiv.org/html/2402.02452v2>
175. Explainable AI Frameworks for Trustworthy Autonomous Cyber ... <https://www.preprints.org/manuscript/202601.1922>
176. [PDF] AI-Based Deepfake Verification Protocols for Legal Evidence <https://www.atlantispress.com/article/126023926.pdf>
177. Summary Plot for CatBoost MultiClass Shap Values #750 - GitHub <https://github.com/shap/shap/issues/750>
178. Catboost tutorial — SHAP latest documentation https://shap.readthedocs.io/en/latest/example_notebooks/tabular_examples/tree_based_models/Catboost%20tutorial.html
179. Multiclass classification using CatBoost and SHAP in Python <https://medium.com/@lochie.links/multiclass-classification-using-catboost-and-shap-in-python-8f2f94d2202a>
180. CatBoost for big data: an interdisciplinary review - PMC <https://pmc.ncbi.nlm.nih.gov/articles/PMC7610170/>

181. How to Combine Scikit-learn, CatBoost, and SHAP for ... <https://machinelearningmastery.com/how-to-combine-scikit-learn-catboost-and-shap-for-explainable-tree-models/>
182. [PDF] Does Class Imbalance Undermine SHAP Explanations? A Multi ... https://assets-eu.researchsquare.com/files/rs-9489204/v1_covered_11b6cc4a-6bf1-4833-9af3-0852395e440e.pdf
183. ShapValues - CatBoost <https://catboost.ai/docs/en/concepts/shap-values>
184. Explaining Multi-class XGBoost Models with SHAP <https://evgenypogorelov.com/multiclass-xgb-shap.html>
185. SHAP based interpretable machine learning for predicting ... <https://www.sciencedirect.com/science/article/pii/S2590123026020116>
186. A Comparative Study of CatBoost and Double Random Forest for Multi ... https://www.researchgate.net/publication/368257066_A_Comparative_Study_of_CatBoost_and_Double_Random_Forest_for_Multi-class_Classification
187. VT Data - H3 Hexagonal Spatial Index Level 9 - Vermont <https://geodata.vermont.gov/datasets/vt-data-h3-hexagonal-spatial-index-level-9-vermont/about>
188. What is a Spatial Index? (H3 Explained) - YouTube https://www.youtube.com/watch?v=_ICYw9vg4ps
189. What is H3 and how does it work in geospatial analysis - Felt <https://felt.com/blog/what-is-h3>
190. Indexing - H3 <https://h3geo.org/docs/highlights/indexing/>
191. Uber H3 Hexagonal Modeling | HEAVY.AI Documentation <https://docs.nvidia.com/heavyai/sql/data-manipulation-dml/geospatial-capabilities/uber-h3-hexagonal-modeling>
192. California wildfire spread derived using VIIRS satellite observations ... <https://pmc.ncbi.nlm.nih.gov/articles/PMC9151742/>
193. Real-Time Wildfire Detection Algorithm Based on VIIRS Fire ... <https://www.mdpi.com/2072-4292/15/6/1541>
194. Satellite (VIIRS) Thermal Hotspots and Fire Activity - ArcGIS Online <https://www.arcgis.com/home/item.html?id=dece90af1a0242dcbf0ca36d30276aa3>
195. [PDF] Assessing the performance of MODIS and VIIRS active fire products in the ... <https://iforest.sisef.org/pdf/?id=ifor3754-015>
196. Wildfire Detection and Vegetation Change Mapping Using VIIRS ... <https://www.youtube.com/watch?v=DKZzXLoiDpE>

197. Comparison of fire radiative power estimates from VIIRS and MODIS ... https://www.researchgate.net/publication/324566694_Comparison_of_fire_radiative_power_estimates_from_VIIRS_and_MODIS_observations
198. [PDF] arXiv:2308.02508v1 [eess.IV] 27 Jul 2023 <https://arxiv.org/pdf/2308.02508>
199. Artificial Intelligence for Early Wildfire Detection in Northern Morocco: A ... <https://doi.org/10.1109/ACCESS.2026.3688023>
200. [PDF] isprs-archives-XLIII-B3-2020-1031 ... <https://isprs-archives.copernicus.org/articles/XLIII-B3-2020/1031/2020/isprs-archives-XLIII-B3-2020-1031-2020.pdf>
201. CatBoost: unbiased boosting with categorical features <https://arxiv.org/abs/1706.09516>
202. Reference papers <https://catboost.ai/docs/en/concepts/educational-materials-papers>
203. CatBoost: paper at NeurIPS 2018 <https://research.yandex.com/blog/catboost-paper-at-neurips-2018>
204. Reviews: CatBoost: unbiased boosting with categorical ... <https://proceedings.neurips.cc/paper/2018/file/14491b756b3a51daac41c24863285549-Reviews.html>
205. CatBoost: gradient boosting with categorical features support https://www.researchgate.net/publication/328576065_CatBoost_gradient_boosting_with_categorical_features_support
206. Prokhorenkova et al (2018) CatBoost: Unbiased boosting with ... <https://www.adrian.idv.hk/2023-05-19-pgvdg18-catboost/>
207. CatBoost's Ordered Boosting and Symmetric Trees <https://apxml.com/courses/getting-started-with-gradient-boosting-algorithms/chapter-5-advanced-gradient-boosting-lightgbm-catboost/catboost-ordered-boosting>
208. CatBoost: unbiased boosting with categorical features <https://scispace.com/papers/catboost-unbiased-boosting-with-categorical-features-5957aimdxs>
209. Prokhorenkova, L., Gusev, G., Vorobev, A ... <https://www.scirp.org/reference/referencespapers?referenceid>
210. CatBoost for big data: an interdisciplinary review <https://link.springer.com/article/10.1186/s40537-020-00369-8>
211. (PDF) CatBoost for big data: an interdisciplinary review https://www.researchgate.net/publication/346496325_CatBoost_for_big_data_an_interdisciplinary_review
212. Effective Methods of Categorical Data Encoding for ... <https://www.mdpi.com/2227-7390/12/16/2553>

213. Accurate predictions on small data with a tabular ... <https://www.nature.com/articles/s41586-024-08328-6>
214. CatBoost: unbiased boosting with categorical features <https://arxiv.org/pdf/1706.09516>
215. Machine Learning for Sensor Analytics - PMC - NIH <https://pmc.ncbi.nlm.nih.gov/articles/PMC12694449/>
216. Gallstone disease classification using SLOA-optimized ... <https://journals.plos.org/plosone/article?id>
217. Cross-Table Pretraining towards a Universal Function ... <https://openreview.net/forum?id>
218. Assessment of environmental and socioeconomic drivers of ... <https://pmc.ncbi.nlm.nih.gov/articles/PMC11845695/>
219. From Local to Global: How One Model Can Forecast ... <https://medium.com/@injure21/from-local-to-global-how-one-model-can-forecast-thousands-of-time-series-a2ac23498902>
220. Recent progress in the data-driven discovery of novel ... <https://www.oaepublish.com/articles/jmi.2022.07>
221. Climate Change AI Workshop Papers <https://www.climatechange.ai/papers>
222. Applied Sciences | December-2 2025 - Browse Articles <https://www.mdpi.com/2076-3417/15/24>
223. peggy1502/Amazing-Resources: List of references and ... <https://github.com/peggy1502/Amazing-Resources>
224. Python 科学技術関連のパッケージ一覧 <https://www.trifields.jp/pypi-science-and-technology-2-6054>
225. Coding Acharya <https://codingacharya.blogspot.com/>
226. docs-merge-05 - 绝不原创的飞龙 <https://www.cnblogs.com/apacheecn/p/18662753>
227. Techniques for deep learning with satellite & aerial imagery <https://github.com/satellite-image-deep-learning/techniques>
228. Spatiotemporal transformer modeling of satellite fire ... <https://www.nature.com/articles/s41598-026-50422-4>
229. Earth Engine Data Catalog <https://developers.google.com/earth-engine/datasets/catalog>
230. Towards Real-Time Earth System Intelligence through ... <https://www.techrxiv.org/doi/pdf/10.36227/techrxiv.177212479.91567605/v1>
231. Advancing horizons in remote sensing: a comprehensive ... <https://link.springer.com/article/10.1007/s00521-024-10165-7>

232. Nighttime light remote sensing for urban applications <https://www.sciencedirect.com/science/article/pii/S0924271623001521>
233. A Multimodal Ensemble Deep Learning Model for Wildfire ... <https://www.mdpi.com/2072-4292/17/19/3310>
234. A temporally consistent global 500 m-resolution monthly VIIRS ... <https://essd.copernicus.org/articles/18/3449/2026/>
235. A Comparative Study of SHAP and LIME in Intrusion Detection Models
CAESYQHrOzAVPrLu4fA70KZCYf3k3pISDxPMU1h615qo8a05DIWylza2mnVPEhcAvkbb
CkD5Zxbl-_cTZS-Okkz1v4yYlKseAcpTKUWig4bQ3IyKU9_vxeQ7jYEQKVsyUpDEAnQ
236. Explainable AI for Digital Forensics: Ensuring Transparency in Legal ...
CAESeQHrOzAVv4JcZnJSIsWyPIuefBvFWcvQvMgRd4xnnZg9NtniMsK6nBtlaHQCNCW
vEX02bfyvIdGVfDYSUzsVyd1evcH9K4CSol8bkclCJbeA0trF3G-
po3f2OQUddXmcTbQ644EDdKjLc_pQXk0gyPdZXXfu33SNNvg
237. (PDF) Developing an Explainable AI System for Digital Forensics
CAES7AEB6zswFShxx2mEx0revd0Ofv3L7wQWW5VReG0YmrXg6T-RVmKcLMUD-
E5IUibYyivyMIoc9lHdVESqchkDmPN5rtOW8f22_eWkTlkOdBKjC_fsTFTJTfpcGzo6Qjjiq
7Fizif7ZVR74sVxRrV_sHHqjt_l-FkAwix6HN2-wPt-LGMRK9Uv2-
_VNc7RDoh_fkdbcxNEAnaYs4RIxYLxdwQsrNPK8TZ-
hMeWiwtnX4NiIBbAPWJGw_LZz8X279qJCvV2WY4aR2VxKlDb73nO8BARZdZpyMKG
bUVXdE4X7f81m2OJPcn1YnOIXGjgeCw
238. Explainable AI for forensic speech authentication within cognitive and ...
CAESkgEB6zswFS9BzezRjlUr5Bt8m-irXPJ90V6UuvN6IWaI6fGC8sNfhCJtzlH-c-
jAnSsWEnTbLm8PqVGiZtYPOp9leBbp8RD-
JhbXGrLOoiFRGxFiABHvjfYNYsuRkLHrjwZSwEIKnwH1HseDsS7v5No8iuDxi4Wr9k1gF
9wi4UBJNQ2JthqgT_44oSa_9MM9DYSnqQ
239. Artificial Intelligence Ethics Framework for the Intelligence Community
CAESawHrOzAVRKZsnSZ4yeyQyqC6KYrNt54AoS_HmF2sp0-
EV2RMLJpSXARD1wJn8HUUT9qFGwDAQlEd87dPeV30YAyPATRIL8aSvb9F-3S9B8nM
QJU9Nji99zi6gsiHUbKLn7A43OpM9si0n7ZZ
240. NIST Publishes Proposed Principles for “Explainable” AI Systems
CAESrwEB6zswFW1VZfoQUZEua2wq1uix__UqN258shukBLZ_w6XYuoTL4UfiJy4GVt7t
5jn1x5xDoBswVDNQpyXBnvZob86fxsKKbwAR2CJQIQlQZEnO7eNSDdJduiO7Y_leZ9z
8czYT_vTfnyvnA-
bRTl09bYtWGiBTl6auLgDiyt2iuOTuQ5S5en_VhgB5p52mHFnhbwYtcpZiETv7lBke3ebB
383zQIC7eGzsba5Chvi3
241. Explainable AI Governance: Frameworks for Trust - Alation
CAESdwHrOzAVLqYOOwun02OGzVEaKNhviiqLhSsP24jN_Xgr-xzb-
uTG7_nogylRhFuk5SYCjc7ZjI-QHca4i7-rg-
qPvxCJG_vH1u_dXVdDAopRkgRci2Ziu5w6gz_EDLmZjZlg7-
iQ1aalvdp2d6UPh9kU_JjaR5ol

242. Artificial intelligence | NIST
CAESZAHrOzAV3MFtuiWb1cbVHIX9B1OXWVCuUUwAzJ1OBQTUdU_M4uLJk5z8b3O
aDcOAqrWvdFicjNbumuusNJgJG39pbw046vrGStVFgL4I5gBc0orgHkQ3Z4Hezivo750Y
ScglQ9A
243. [PDF] Explainable Artificial Intelligence (XAI) Models for Transparent ... - ijsrm
CAEScAHrOzAVfF-Gh-
yi7k2s8Rh74bi6IbRa_VkSOiVsIpLB3jToHT2DGYPl802CZrctMStCw_h5DBVd-
t0mx_0dhUMqvwH4IddCyvoaBjwKpQ1sXirQONv0Wx2jK4bHVNf_m7qjvwMIft0HZ-3g
WQw_Thc
244. A Cryptographic Verification Standard for Agentic AI Governance in ... - IETF
CAESjQEB6zswFelIEIbPXrWiSXWwSww49051yfggbnLl4WW9DQI7cw0EtHFHWdMx3
C0PGqtbNRPhAchZg5mZRtbTnBJf9kOZ95KAxtJ8CExJZedp6L--
qej6Le0stZ4XlernPuaFDtvRdphDl_1_4TNG2JHIVOOyt5ZDFIo4kVfobBfjhnqVyLAmbk0
4RDY5_E8
245. shap.TreeExplainer — SHAP latest documentation <https://shap.readthedocs.io/en/latest/generated/shap.TreeExplainer.html>
246. A Gentle Introduction to SHAP for Tree-Based Models <https://machinelearningmastery.com/a-gentle-introduction-to-shap-for-tree-based-models/>
247. Understanding SHAP for multi-classification problem #367 <https://github.com/shap/shap/issues/367>
248. Does Class Imbalance Undermine SHAP Explanations? A ... https://www.researchgate.net/publication/404129962_Does_Class_Imbalance_Undermine_SHAP_Explanations_A_Multi-Domain_Empirical_Study_of_Feature_Attribution_Stability
249. How to get SHAP values for each class on a multiclass ... <https://stackoverflow.com/questions/71753428/how-to-get-shap-values-for-each-class-on-a-multiclass-classification-problem-in>
250. A comparative analysis of machine learning models in ... <https://arxiv.org/html/2604.07258v1>
251. TreeSHAP: Optimized for Tree-Based Models <https://apxml.com/courses/model-interpretability-explainability/chapter-3-shap-additive-explanations/treeshap-explainer>
252. Harnessing the Power of H3-Py: A Practitioner's Guide to ... <https://medium.com/aimonks/harnessing-the-power-of-h3-py-a-practitioners-guide-to-hexagonal-spatial-indexing-108ded50fb3b>
253. Discrete Global Grid Systems as scalable geospatial ... - PMC <https://pmc.ncbi.nlm.nih.gov/articles/PMC8958999/>
254. Solving Geospatial Analytics Performance Bottleneck <https://www.e6data.com/blog/geospatial-analytics-performance-bottleneck-h3-vs-quadkey-for-spatial-indexing>

255. Gradient Boosting Revisited: Comparative Analysis of ... <https://www.sciencepublishinggroup.com/article/10.11648/j.mlr.20261101.14>
256. jxucoder/awesome-gradient-boosting-machines <https://github.com/jxucoder/awesome-gradient-boosting-machines>
257. Comparison of Gradient Boosting Decision Tree Algorithms ... https://www.researchgate.net/publication/351133481_Comparison_of_Gradient_Boosting_Decision_Tree_Algorithms_for_CPU_Performance
258. Satellite (VIIRS) Thermal Hotspots and Fire Activity - ArcGIS Hub https://hub.arcgis.com/datasets/dece90af1a0242dcbf0ca36d30276aa3_0/about
259. Improving Nocturnal Fire Detection With the VIIRS Day– ... https://www.researchgate.net/publication/304369177_Improving_Nocturnal_Fire_Detection_With_the_VIIRS_Day-Night_Band
260. Fire Hotspots Mapping and Forecasting in Indonesia Using ... https://repository.uir.ac.id/23068/1/Fire_Hotspots_Mapping_and_Forecasting_in_Indonesia_Using_Deep_Learning_Algorithm-3.pdf
261. Automated forest fire detection in ecological monitoring using ... <https://pmc.ncbi.nlm.nih.gov/articles/PMC12808325/>
262. Deep Learning Approach to Improve Spatial Resolution ... <https://scholarwolf.unr.edu/entities/thesis/e5a1e20a-7c66-4371-9456-d0210127fdb>
263. (PDF) CatBoost-SHAP analysis decodes spatial synergies ... https://www.researchgate.net/publication/403857457_CatBoost-SHAP_analysis_decodes_spatial_synergies_driving_prehistoric_architectural_elongation_on_the_Loess_Plateau
264. An Explainable CatBoost Model for Crater Classification ... <https://www.mdpi.com/2072-4292/17/7/1236>
265. Hybrid bagging and boosting with SHAP based feature ... - PMC <https://pmc.ncbi.nlm.nih.gov/articles/PMC11652632/>
266. Data Leakage in Applying Geospatial Foundation Models <https://research.ibm.com/publications/data-leakage-in-applying-geospatial-foundation-models>
267. SHAP for Categorical Features with CatBoost <https://medium.com/data-science/shap-for-categorical-features-with-catboost-8315e14dac1>
268. shap_values_tutorial.ipynb https://github.com/catboost/catboost/blob/master/catboost/tutorials/model_analysis/shap_values_tutorial.ipynb
269. CatBoost for Big Data: an Interdisciplinary Review https://assets-eu.researchsquare.com/files/rs-54646/v2_covered.pdf

270. Comparison of the CatBoost Classifier with other Machine ... https://www.researchgate.net/publication/347351801_Comparison_of_the_CatBoost_Classifier_with_other_Machine_Learning_Methods
271. VIIRS Nightfire Super-Resolution Method for Multiyear ... <https://www.mdpi.com/2072-4292/18/2/314>
272. A comprehensive survey of the machine learning pipeline ... <https://www.sciencedirect.com/science/article/pii/S1574954125003346>
273. A Transfer Learning Study From Landsat-8 to Sentinel-2 <http://ieeexplore.ieee.org/iel8/4609443/10330207/10620606.pdf>
274. Daily R Archives <https://dailyr.netlify.app/archives/>
275. A Report on CatBoost: unbiased boosting with categorical ... https://www.researchgate.net/publication/389249568_A_Report_on_CatBoost_unbiased_boosting_with_categorical_features
276. [PDF] CatBoost: unbiased boosting with categorical features <https://www.semanticscholar.org/paper/CatBoost%3A-unbiased-boosting-with-categorical-Ostroumova-Gusev/ee0a0f04d45f86bf50b24d7258e884725fcaa621>
277. MonoForest framework for tree ensemble analysis https://proceedings.neurips.cc/paper_files/paper/2019/file/1b9a80606d74d3da6db2f1274557e644-Paper.pdf
278. Machine learning approaches for assessing stability in acid ... <https://www.sciencedirect.com/science/article/pii/S1995822624002528>
279. CatBoost: gradient boosting with categorical features support <https://arxiv.org/abs/1810.11363>
280. Full article: Global and local magnitude and spatial pattern ... <https://www.tandfonline.com/doi/full/10.1080/15481603.2023.2297549>
281. CatBoost: unbiased boosting with categorical features <https://scispace.com/pdf/catboost-unbiased-boosting-with-categorical-features-5957aimdxs.pdf>
282. A Comparative Analysis of Gradient Boosting Algorithms <https://repositorio.uam.es/bitstreams/61b6665a-5914-4559-ad8e-13d069cc0817/download>
283. Elasticsearch for big geotemporal data | Zhyrenkov <https://pp.isoftware.kiev.ua/index.php/ojs1/article/view/764>
284. A Benchmark Dataset and Cross-Modality Attention Model <https://arxiv.org/abs/2410.15689>
285. APA <https://www.jport.co/index.php/jport/citationstylelanguage/get/apa?submissionId>
286. TreeExplainer support for GradientBoostingClassifier in ... <https://github.com/shap/shap/issues/2887>

287. Comparative Analysis of Tree-Based Ensemble Learning ... <https://www.mdpi.com/2073-4441/15/14/2661>
288. a global terrorism data-driven comparison of machine learning ... <https://link.springer.com/article/10.1007/s13278-025-01498-9>
289. SHAP values with examples applied to a multi- ... <https://harpomaxx.github.io/post/shap-values/>
290. Artificial intelligence <https://www.nist.gov/artificial-intelligence>
291. Explainable AI for Comparative Analysis of Intrusion ... <https://arxiv.org/html/2406.09684v1>
292. Explainable AI for cyber security. Improving transparency ... https://ijaem.net/issue_dcp/Explainable%20AI%20for%20cyber%20security.%20Improving%20transparency%20and%20trust%20in%20intrusion%20detection%20systems.pdf
293. Harnessing Artificial Intelligence to Fulfill National Security ... <https://www.intelligence.gov/ai>
294. NIST AI Agent Standards: What It Means for Enterprise ... <https://labs.cloudsecurityalliance.org/research/csa-research-note-nist-ai-agent-standards-initiative-2026040/>
295. Understanding AI's Capabilities: NIST Launches Program ... <https://www.ansi.org/standards-news/all-news/5-31-24-understanding-ai-capabilities-nist-launches-program-to-advance-sociotechnical-testing>
296. Why AI Security Will Define the Future of Trust <https://www.cfr.org/articles/securing-intelligence-why-ai-security-will-define-future-trust>
297. National Institute of Standards and Technology (NIST) ... https://downloads.regulations.gov/NIST-2023-0009-0035/attachment_1.pdf
298. Engineering Spatial Features for Machine Learning Success https://www.linkedin.com/posts/mbforr_spatial-data-is-a-goldmine-but-how-you-engineer-activity-7433343904109940736-IHWx
299. 5 Spatial ML Mistakes to Avoid in Nigeria https://www.linkedin.com/posts/spatial-and-data-science-society_spatialdatascience-machinelearning-geospatialai-activity-7432753635656146944-pdQq
300. PyPI Download Leaderboard <https://pypilb.vercel.app/>
301. タグ一覧(アルファベット順)【直近1年間/上位 ... <https://qiita.com/j5c8k6m8/items/6b15a81204a1d458e392>
302. Wildfire Susceptibility Mapping Using Deep Learning and ... <https://www.mdpi.com/2571-6255/8/10/407>

303. Artificial intelligence and remote sensing frameworks for ... <https://www.aimspress.com/article/id/698d5e2cba35de4e3973690c>
304. UNFCCC Focal Point of the Philippines https://redd.unfccc.int/files/philippine_frl_document_final_6dec2022.pdf
305. Leveraging the Internet of Things, Remote Sensing, and ... https://www.researchgate.net/publication/388122231_Leveraging_the_Internet_of_Things_Remote_Sensing_and_Artificial_Intelligence_for_Sustainable_Forest_Management
306. Bridging Climate and Socio-Environmental Vulnerability for ... <https://www.mdpi.com/1999-4907/17/2/182>
307. Cultivating historical heritage area vitality using urban ... <https://eprints.gla.ac.uk/255489/1/255489.pdf>
308. Korean Journal of Anesthesiology <https://ekja.org/journal/CitedBy.php?number>
309. Advancements in Artificial Intelligence Applications for ... <https://www.mdpi.com/1999-4907/16/4/704>
310. Effective spatial embeddings for tabular data <https://mlumiste.com/technical/spatial-embeddings/>
311. An Optimized CatBoost Model for Spatiotemporal ... <https://www.mdpi.com/2072-4292/18/15/2579>
312. The 37 Geospatial Python Packages You Definitely Need <https://forrest.nyc/the-37-geospatial-python-packages-you-definitely-need/>
313. Edge-intelligent leak detection in water distribution systems ... https://www.researchgate.net/publication/397015283_Edge-intelligent_leak_detection_in_water_distribution_systems_using_CatBoost_A_sustainable_solution_for_reducing_infrastructure_losses
314. IOAH3: Importance-Driven Adaptive Spatial Partitioning <https://arxiv.org/pdf/2606.18280>
315. Hybrid Spatial-Temporal Deep Learning for Privacy ... <https://www.techrxiv.org/doi/10.36227/techrxiv.176799976.66603504>
316. Deep learning for fog detection and prediction <https://www.sciencedirect.com/science/article/pii/S2590005626002304>
317. A review of machine learning applications in wildfire ... <https://cdnsiencepub.com/doi/10.1139/er-2020-0019>
318. Predicting household piped water consumption using satellite ... <https://journals.plos.org/water/article?id>
319. Deep transfer learning and attention based P2.5 forecasting in ... <https://pmc.ncbi.nlm.nih.gov/articles/PMC12394441/>

320. FlameGuard: an AI-driven smart map for early detection ... <https://www.nature.com/articles/s41598-025-28814-9>
321. Remote Sensing for Wildfire Mapping: A Comprehensive ... <https://www.mdpi.com/2571-6255/8/8/316>
322. Hyperspectral remote sensing of fire https://research.vu.nl/files/249559292/Hyperspectral_remote_sensing_of_fire_State_of_the_art_and_future_perspectives.pdf
323. CatBoost Part 1: Ordered Target Encoding - YouTube <https://www.youtube.com/watch?v=KXOTSkPL2X4>
324. Explainable artificial intelligence models using SHAP enhanced ... <https://pmc.ncbi.nlm.nih.gov/articles/PMC13332054/>
325. Data Products: Fire/Hot Spot Characterization | GOES-R Series <https://www.goes-r.gov/products/baseline-fire-hot-spot.html>
326. CatBoost: Mastering Target Encoding and Ordered Boosting - Towards AI <https://pub.towardsai.net/catboost-mastering-target-encoding-and-ordered-boosting-b018bec09757>
327. CatBoost Ordered Target Statistics - ApX Machine Learning <https://apxml.com/courses/mastering-gradient-boosting-algorithms/chapter-6-catboost-gradient-boosting/catboost-ordered-ts>
328. Improving Multi-Class Classification for Recognition of the ... <https://www.mdpi.com/2076-3417/15/13/7071>
329. Selecting and Interpreting Multiclass Loss and Accuracy ... https://www.researchgate.net/publication/377833097_Selecting_and_Interpreting_Multiclass_Loss_and_Accuracy_Assessment_Metrics_for_Classifications_with_Class_Imbalance_Guidance_and_Best_Practices
330. Multi Class Classification - Arize AX Docs <https://arize.com/docs/ax/machine-learning/machine-learning/use-cases-ml/multiclass-classification>
331. Explainable AI in Clinical Decision Support Systems - PMC - NIH <https://pmc.ncbi.nlm.nih.gov/articles/PMC12427955/>
332. How Reliable and Stable are Explanations of XAI Methods? <https://arxiv.org/html/2407.03108v1>
333. Explainable Multi-class Classification of Medical Data <https://www.semanticscholar.org/paper/Explainable-Multi-class-Classification-of-Medical-Hu-Sokolova/296a6e321f4f19f1fa8ef0a78a14d4a521955d67>
334. JOURNAL OF INFORMATION SYSTEMS APPLIED ... <https://jisara.org/2026-19/n1/JISARAv19n1.pdf>
335. Publications <https://krlgroup.github.io/publications/>
336. Publications — ATD AI <https://www.atdai.ai/publications>

337. [R] Our new classification algorithm outperforms CatBoost, XGBoost ... https://www.reddit.com/r/MachineLearning/comments/1cq5y4/r_our_new_classification_algorithm_outperforms/
338. Why CatBoost is the hidden gem of tabular AI (and ... - Valeriy's Substack <https://valeman.substack.com/p/why-catboost-is-the-hidden-gem-of>
339. XGBoost vs. CatBoost: Explained Visually - YouTube <https://www.youtube.com/watch?v=Qr-cbaNecoE>
340. Lightgbm vs xgboost vs catboost - Data Science Stack Exchange <https://datascience.stackexchange.com/questions/49567/lightgbm-vs-xgboost-vs-catboost>
341. XGBoost vs. LightGBM vs. CatBoost: A Comparison - ApX Machine Learning <https://apxml.com/courses/getting-started-with-gradient-boosting-algorithms/chapter-5-advanced-gradient-boosting-lightgbm-catboost/xgboost-vs-lightgbm-vs-catboost>
342. idea behind xgboost/lightgbm/catboost in ... <https://stackoverflow.com/questions/53916535/idea-behind-xgboost-lightgbm-catboost-in-comparison>
343. Explainable AI in Cyber Defense: Why Transparency ... <https://medium.com/@gupta.brij/explainable-ai-in-cyber-defense-why-transparency-matters-more-than-ever-c1b252f0eba0>
344. A systematic review on the integration of explainable artificial ... <https://pmc.ncbi.nlm.nih.gov/articles/PMC11877648/>
345. Use of Explainable Artificial Intelligence for Analyzing and ... <https://www.mdpi.com/2073-431X/14/5/160>
346. Explainable Artificial Intelligence (XAI) for Trustworthy ... https://www.researchgate.net/publication/396907619_Explainable_Artificial_Intelligence_XAI_for_Trustworthy_Cybersecurity_Systems_Investigating_Transparency_and_Interpretability_of_AI_Models_in_Cyber_Defense_Decision-Making
347. High-Risk AI Models Need Military-Grade Security <https://warontherocks.com/cogs-of-war/high-risk-ai-models-need-military-grade-security/>
348. Explainable AI in National Security: Enhancing Trust and ... <https://www.jetir.org/papers/JETIR2312786.pdf>
349. Satellite (VIIRS) Thermal Hotspots and Fire Activity <https://climate.esri.ca/search?tags=viirs>
350. Smoothing target encoding and class center-based firefly algorithm for ... <https://journalofbigdata.springeropen.com/articles/10.1186/s40537-022-00679-z>
351. Satellite (VIIRS) Thermal Hotspots and Fire Activity - National Data Platform <https://nationaldataplatfrom.org/catalog/dataset/satellite-viirs-thermal-hotspots-and-fire-activity>

352. Datasets tagged viirs in Earth Engine - Google for Developers <https://developers.google.com/earth-engine/datasets/tags/viirs>
353. Downsides of target encoding compared to one hot ... - Reddit https://www.reddit.com/r/datascience/comments/1ayhbx/b/ Downsides_of_target_encoding_compared_to_one_hot/
354. Encoding Categorical Variables: A Deep Dive into Target ... - Medium <https://medium.com/data-science/encoding-categorical-variables-a-deep-dive-into-target-encoding-2862217c2753>
355. Target Encoding and the Lost Art of Feature Engineering - LinkedIn <https://www.linkedin.com/pulse/target-encoding-lost-art-feature-engineering-daniele-micci-barreca>
356. CatBoost, LightGBM, XGBoost: CPU vs GPU <https://www.kaggle.com/code/marychin/catboost-lightgbm-xgboost-cpu-vs-gpu/script?scriptVersionId>
357. Comparative study of the 3 most commonly used boosting ... <https://medium.com/@hamzamlwh/comparative-study-of-the-3-most-commonly-used-boosting-methods-a49b1863e781>
358. 1st Place Solution - CatBoost All The Way Down <https://www.kaggle.com/competitions/playground-series-s4e10/writeups/hardy-xu-1st-place-solution-catboost-all-the-way-d>
359. Smarter Geospatial Predictions with CatBoost: A Powerful ... <https://medium.com/@altair0622/smarter-geospatial-predictions-with-catboost-a-powerful-ally-for-gis-practitioners-855b6586c83b>
360. Decentralized geoprivacy: leveraging social trust on the ... <https://search.dataone.org/view/sha256%3Acab70d23f4d8ec34b2d741b2ce0fb26d061e130640b9dbb455485cc9338cc101>
361. CatBoost Classifier Optimized by the Lyrebird Algorithm to ... <https://www.mdpi.com/1999-5903/16/10/381>
362. https://scholar.google.com/citations?view_op=view_... https://scholar.google.com/citations?view_op=view_citation&hl=en&user=2hVUXWMAAAAJ&citation_for_view=2hVUXWMAAAAJ:dshw04ExmUIC
363. OBPS Repository - Ocean Best Practices System <https://repository.oceanbestpractices.org/handle/11329/1103>
364. Wildfire Detection Using Vision Transformer Models <https://peer.asee.org/54966.pdf>
365. Utilizing MODIS remote sensing and integrated data for ... <https://www.cabidigitallibrary.org/doi/full/10.5555/20250326236>

366. Science Team - MODIS Web - NASA https://modis.gsfc.nasa.gov/sci_team/pubs/citations_new.php?year
367. Multi-source remote sensing-based forest fire monitoring <https://isprs-archives.copernicus.org/articles/XLVIII-M-5-2024/189/2025/isprs-archives-XLVIII-M-5-2024-189-2025.pdf>
368. Share some insights about LightGBM vs. CatBoost in Zillow ... <https://www.kaggle.com/competitions/zillow-prize-1/writeups/mlin-share-some-insights-about-lightgbm-vs-catboos>
369. Is CatBoost faster than LightGBM and XGBoost? <https://deliveryhero.jobs/blog/is-catboost-faster-than-lightgbm-and-xgboost/>
370. Decission between boosting algorithms like XGBoost ... <https://www.kaggle.com/questions-and-answers/550926>
371. Closing the AI Audit Gap: Forensic Design Requirements ... <https://papers.ssrn.com/sol3/Delivery.cfm/7101678.pdf?abstractid>
372. Adversarial Machine Learning: A Taxonomy and ... <https://nvlpubs.nist.gov/nistpubs/ai/NIST.AI.100-2e2025.pdf>
373. Admissibility of AI-Generated Forensic Evidence <https://zenodo.org/records/17500004>
374. Investing AI Ethics in Forensic Investigations <https://www.athensjournals.gr/business/2025-11-4-2-Coltri.pdf>
375. Responsible AI in Financial Crime Compliance: A Path ... <https://www.amlrightsourc.com/resources/responsible-ai-in-financial-crime-compliance-a-path-towards-effective-and-ethical-ai-assisted-anti-money-laundering-and-counter-terrorist-financing>
376. The Ethics of AI in Criminal Justice: Balancing Innovation ... <https://tranquility.ai/formatted-articles/the-ethics-of-ai-in-criminal-justice-balancing-innovation-and-responsibility>
377. The ethics of facial recognition technologies, surveillance, and ... <https://pmc.ncbi.nlm.nih.gov/articles/PMC8320316/>
378. A Review of TRiSM Frameworks in Artificial Intelligence ... <https://onlinelibrary.wiley.com/doi/10.1111/exsy.70213>
379. CatBoost vs XGBoost vs LightGBM vs scikit-learn ... <https://www.linkedin.com/pulse/boosting-techniques-battle-catboost-vs-xgboost-lightgbm-uttam-kumar>
380. XGBoost vs LightGBM vs CatBoost: The Gradient Boosting ... <https://pristren.com/blog/xgboost-lightgbm-catboost/>
381. (PDF) Fire Prediction Based on CatBoost Algorithm https://www.researchgate.net/publication/353373004_Fire_Prediction_Based_on_CatBoost_Algorithm

382. Enhancing Data Security in Satellite Communication Systems <https://inspirehep.net/files/0776e84a48bef31143b5dead999cfee6>
383. VIIRS Hotspot - Daily - ASMC - Asean <https://asmc.asean.org/asmc-haze-hotspot-daily-new/>
384. Comparing Satellite Data for Next-Day Wildfire Predictability <https://arxiv.org/pdf/2503.08580>
385. GA-CatBoost-Weight Algorithm for Predicting Casualties in ... <https://www.mdpi.com/2227-7390/12/6/818>
386. Auto ML Tabular When reading the paper <https://arxiv.org/> ... <https://www.facebook.com/groups/MDLI1/posts/1805583872938986/>
387. Persistent homology for MCI classification: a comparative ... <https://www.frontiersin.org/journals/neuroscience/articles/10.3389/fnins.2025.1518984/full>
388. (PDF) Boosting methods for multi-class imbalanced data ... https://www.researchgate.net/publication/344040059_Boosting_methods_for_multi-class_imbalanced_data_classification_an_experimental_review
389. Improving GBDT Performance on Imbalanced Datasets <https://arxiv.org/html/2407.14381v1>
390. XGBoost & LightGBM & Catboost - Imbalanced Data <https://www.kaggle.com/code/kaanboke/xgboost-lightgbm-catboost-imbalanced-data>
391. (PDF) Boosting methods for multi-class imbalanced data ... <https://scispace.com/papers/boosting-methods-for-multi-class-imbalanced-data-2b670a549f>
392. Handling Imbalanced Datasets with XGBoost <https://medium.com/@mate.voros1998/handling-imbalanced-datasets-with-xgboost-optimizing-model-performance-with-smart-parameter-tuning-18568c7783cf>
393. A Comparative Study of XAI Methods on Imbalanced Data <https://papers.ssrn.com/sol3/Delivery.cfm/6068409.pdf?abstractid>
394. XGBoost for multiclassification and imbalanced data <https://stackoverflow.com/questions/67868420/xgboost-for-multiclassification-and-imbalanced-data>
395. How to Configure XGBoost for Imbalanced Classification <https://www.machinelearningmastery.com/xgboost-for-imbalanced-classification/>
396. VIIRS I-Band 375 m Active Fire Data <https://www.earthdata.nasa.gov/data/instruments/viirs/viirs-i-band-375-m-active-fire-data>
397. NIST PRINCIPLES FOR EXPLAINABLE AI - STIP Compass <https://stip.oecd.org/stip/interactive-dashboards/policy-initiatives/2023%2Fdata%2FpolicyInitiatives%2F26746>
398. Artificial Intelligence Risk Management Framework <https://www.regulations.gov/document/NIST-2021-0004-0001>

399. Unpacking New NIST Guidance on Artificial Intelligence <https://techpolicy.press/unpacking-new-nist-guidance-on-artificial-intelligence>
400. NIST adds to AI security guidance with Cybersecurity ... <https://www.cybersecuritydive.com/news/nist-ai-cybersecurity-framework-profile/808134/>
401. NIST Releases Risk Management Framework for Building ... <https://www.executivegov.com/articles/nist-releases-risk-management-framework-for-building-trustworthy-ai-systems>
402. NIST's AI Risk Management Framework: What It Is, Why It ... <https://blacksmithinfosec.com/nists-ai-risk-management-framework-what-it-is-why-it-exists-and-what-msps-should-know/>
403. Stability of SHAP-Based Feature Importance Ranking ... <https://journal.unugiri.ac.id/index.php/statkom/article/view/6455>
404. An example of SHAP and XGBoost <https://eprints.gla.ac.uk/273108/1/273108.pdf>
405. CatBoost-LightGBM-XGBoost Explained by SHAP <https://www.kaggle.com/code/kaanboke/catboost-lightgbm-xgboost-explained-by-shap>
406. A Comparative Study of XGBoost and CatBoost <http://www.diva-portal.org/smash/record.jsf?pid>
407. Active Fire Detection <https://forest-fire.emergency.copernicus.eu/about-effis/technical-background/active-fire-detection>
408. Advancing wildfire susceptibility mapping through machine ... <https://www.frontiersin.org/journals/forests-and-global-change/articles/10.3389/ffgc.2025.1705341/full>
409. CatBoost: unbiased boosting with categorical features <https://papers.nips.cc/paper/7898-catboost-unbiased-boosting-with-categorical-features.pdf> NeurIPS.Papers+1 NeurIPS.Papers+1
410. Predicting thyroid cancer recurrence using supervised CatBoost <https://pmc.ncbi.nlm.nih.gov/articles/PMC12129522/>
411. A Comprehensive Modeling Framework for Air Quality ... <https://ui.adsabs.harvard.edu/abs/2025PApGe.182.4771A/abstract>
412. ShapValues https://catboost.ai/docs/en/concepts/output-data_feature-analysis_shap-values
413. Optimizing Faulting Prediction for Rigid Pavements Using a ... <https://pdfs.semanticscholar.org/d395/b4844aae856c8cf53f3a0ff84334923b847c.pdf>
414. **XGBoost vs. LightGBM vs. CatBoost** <https://www.kaggle.com/general/713178>
415. Performance of Various Artificial Intelligence Models for ... <https://hal.science/hal-05173646/document>

416. Machine Learning for Improved Threat Detection <https://jcbi.org/index.php/Main/article/view/499>
417. Comparative Study of CatBoost, XGBoost, and LightGBM ... <https://jisis.org/article/2023.I4.001/70886/>
418. Comparison of machine learning model performance for ... <https://pmc.ncbi.nlm.nih.gov/articles/PMC12222744/>
419. SHAP TreeExplainer is not compatible with XGBoost ... <https://github.com/shap/shap/issues/2662>
420. Implementing Responsible Artificial Intelligence in the ... <https://media.defense.gov/2021/May/27/2002730593/-1/-1/0/IMPLEMENTING-RESPONSIBLE-ARTIFICIAL-INTELLIGENCE-IN-THE-DEPARTMENT-OF-DEFENSE.PDF>
421. Anthropic's Advanced AI Framework <https://www-cdn.anthropic.com/files/4zrzovbb/website/0a58d567024a8b448ff15158ebc3625328dfcc1f.pdf>
422. Five Questions the US Government Should Answer About ... <https://www.techpolicy.press/five-questions-the-us-government-should-answer-about-its-secretive-frontier-ai-framework/>
423. Understanding AI Safety: Principles, Frameworks, and Best ... <https://www.tigera.io/learn/guides/llm-security/ai-safety/>
424. Anthropic vs. Washington: AI Ethics Collide with National ... <https://www.jdsupra.com/legalnews/anthropic-vs-washington-ai-ethics-7224675/>
425. Anthropic Refuses Pentagon Demand to Remove AI ... <https://www.asisonline.org/security-management-magazine/latest-news/today-in-security/2026/february/Anthropic-Refusal/>
426. Governance and Ethics in Military AI <https://www.concretecms.com/about/blog/ai/governance-and-ethics-in-military-ai>
427. DoD AI Ethical Principles - STIP Compass <https://stip.oecd.org/stip/interactive-dashboards/policy-initiatives/2025%2Fdata%2FpolicyInitiatives%2F26945>
428. Ethics and regulation of AI in defence technology <https://www.taylorwessing.com/en/interface/2025/defence-tech/ethics-and-regulation-of-ai-in-defence-technology>
429. CatBoost: unbiased boosting with categorical features <https://wiki.math.uwaterloo.ca/statwiki/index.php?title>
430. A taxonomy for detecting and preventing temporal data leakage in ... <https://journals.plos.org/plosone/article?id=10.1371/journal.pone.0340167>
431. Enhancing Maritime Trajectory Forecasting via H3 Index and Causal ... <https://arxiv.org/html/2405.09596v1>
432. H3: Home <https://h3geo.org/>

433. uber/h3: Hexagonal hierarchical geospatial indexing system - GitHub <https://github.com/uber/h3>
434. shap/shap/explainers/_tree.py at master https://github.com/slundberg/shap/blob/master/shap/explainers/_tree.py
435. Application of Random Forest and SHAP Tree Explainer in ... <https://www.mdpi.com/2220-9964/10/9/629>
436. treeshap <https://modeloriented.github.io/treeshap/>
437. Difference in model feature importance and SHAP ... <https://stackoverflow.com/questions/71088769/difference-in-model-feature-importance-and-shap-summary-plot>
438. When to Use "XGBoost | LightGBM | CatBoost" <https://www.kaggle.com/code/masayakawamata/when-to-use-xgboost-lightgbm-catboost>
439. Titanic: Keras vs LightGBM vs CatBoost vs XGBoost <https://www.kaggle.com/code/ajalnine/titanic-keras-vs-lightgbm-vs-catboost-vs-xgboost>
440. Keep AI Testing Defense-Worthy <https://www.lawfaremedia.org/article/keep-ai-testing-defense-worthy>
441. Ethical Norms for New Generation Artificial Intelligence ... <https://cset.georgetown.edu/publication/ethical-norms-for-new-generation-artificial-intelligence-released/>
442. New AI Incident Reporting Program Proposed for DOD <https://www.wiley.law/alert-New-AI-Incident-Reporting-Framework-Proposed-for-DOD>
443. The Three C's To Consider Before Deploying AI ... <https://www.clarkhill.com/news-events/news/the-three-cs-to-consider-before-deploying-ai-technologies-contracts-compliance-and-culture-2/>
444. From Linear Regression to Boosted Trees https://www.researchgate.net/publication/396584864_Comparative_Analysis_of_Machine_Learning_Models_for_House_Price_Prediction_From_Linear_Regression_to_Boosted_Trees
445. Categorical Boosting – PSO: A SHAP – Interpretable Hybrid ... <https://onlinelibrary.wiley.com/doi/full/10.1002/eng2.70979>
446. Prediction of Critical Filling of a Storage Area Network by ... <https://www.mdpi.com/2079-9292/11/24/4150>
447. Techniques for mitigating overfitting in machine learning <https://www.frontiersin.org/journals/artificial-intelligence/articles/10.3389/frai.2026.1794271/pdf>
448. Federated SHAP-Based Intrusion Detection System for IoMT <https://www.scitepress.org/Papers/2026/143155/143155.pdf>
449. Machine Learning Model for Leak Detection Using Water ... <https://pmc.ncbi.nlm.nih.gov/articles/PMC10648304/>

450. From Local Explanations to Global Understanding with ... - PMC <https://pmc.ncbi.nlm.nih.gov/articles/PMC7326367/>
451. Explanation Multiplicity in SHAP <https://arxiv.org/pdf/2601.12654>
452. (PDF) Application of Random Forest and SHAP Tree ... https://www.researchgate.net/publication/354750109_Application_of_Random_Forest_and_SHAP_Tree_Explainer_in_Exploring_Spatial_InJustice_to_Aid_Urban_Planning
453. A systematic review of Explainable Artificial Intelligence ... <https://www.sciencedirect.com/science/article/pii/S277266222300070X>
454. Evidential Hybrid Re-sampling for Multi-class Imbalanced ... <https://hal.science/hal-03764086/document>
455. Evaluating XAI techniques under class imbalance using ... <https://pmc.ncbi.nlm.nih.gov/articles/PMC12657415/>
456. XGBoost vs. LightGBM vs. CatBoost <https://apxml.com/posts/xgboost-vs-lightgbm-vs-catboost>
457. CatBoost vs XGBoost vs LightGBM: Which to Use (2026) <https://mlsimplified.com/gradient-boosting-xgboost-lightgbm-catboost/>
458. S6E4 | XGBoost + LightGBM + CatBoost | Kaggle <https://www.kaggle.com/code/sarveshchhetri/s6e4-xgboost-lightgbm-catboost>
459. XGBoost vs LightGBM vs CatBoost vs AdaBoost | Kaggle <https://www.kaggle.com/faressayah/xgboost-vs-lightgbm-vs-catboost-vs-adaboost/code>
460. Comparison between XGBoost, LightGBM and CatBoost Using a ... <https://www.semanticscholar.org/paper/Comparison-between-XGBoost%2C-LightGBM-and-CatBoost-a-Daoud/b992fdb71b4b78d7b81dc3761402f4eb446077c2>
461. CatBoost, LightGBM, XGBoost: CPU vs GPU | Kaggle <https://www.kaggle.com/marychin/catboost-lightgbm-xgboost-cpu-vs-gpu/execution>
462. Investigation of the impact effectiveness of adversarial data ... https://www.itm-conferences.org/articles/itmconf/pdf/2024/02/itmconf_hmmocs2023_04011.pdf
463. CatBoost Classifier Improved by the Lyrebird Algorithm to ... <https://ijsdr.org/papers/IJSDR2503262.pdf>
464. arXiv:2105.06141v1 [q-bio.QM] 13 May 2021 <https://arxiv.org/pdf/2105.06141>
465. Evaluating the Effectiveness of the CatBoost Classifier in ... <https://www.semanticscholar.org/paper/Evaluating-the-Effectiveness-of-the-CatBoost-in-FTP-Hajjouz-Avksentieva/209f11bdb88328919133d35c898e639b2ee2bd50>
466. CatBoost-Based Intrusion Detection System with NearMiss ... <https://www.ijcna.org/abstract.php?id>

467. NTIA AI Accountability RFC <https://www.regulations.gov/document/NTIA-2023-0005-0001>
468. Safe, Secure, and Trustworthy Development and Use of ... <https://www.federalregister.gov/documents/2023/11/01/2023-24283/safe-secure-and-trustworthy-development-and-use-of-artificial-intelligence>
469. Artificial Intelligence: Statutes, Policies & Guidance (SPG) <https://www.waru.edu/artificial-intelligence/statutes-policies-guidance>
470. AI strategies and compliance plan <https://www.gsa.gov/artificial-intelligence/resources/ai-strategies-and-compliance-plan>
471. AI Regulation & Government Use of AI <https://www.aiethicist.org/ai-regulations>
472. Artificial Intelligence Accountability Policy <https://www.ntia.gov/issues/artificial-intelligence/ai-accountability-policy-report/overview>
473. What We're Reading on AI Regulation <https://cset.georgetown.edu/article/what-were-reading-on-ai-regulation/>
474. Why Model Explanations Break: SHAP, LIME and Tree ... <https://atul4u.medium.com/why-model-explanations-break-shap-lime-and-tree-ensembles-in-practice-4de6ce126342>
475. Explainable artificial intelligence models using SHAP ... https://www.researchgate.net/publication/404393102_Explainable_artificial_intelligence_models_using_SHAP_enhanced_CatBoost_Bi-GRU_with_attention_and_Tab_Transformer
476. Explaining XGBoost predictions with SHAP value <https://journals.vilniustech.lt/index.php/NTCS/article/view/17901>
477. arXiv:2302.08434v1 [cs.LG] 16 Feb 2023 <https://arxiv.org/pdf/2302.08434>
478. 30 Days of ML <https://www.kaggle.com/competitions/30-days-of-ml/discussion/265781>
479. Comparative Study of CatBoost, XGBoost, and LightGBM ... https://www.researchgate.net/publication/376180322_Comparative_Study_of_CatBoost_XGBoost_and_LightGBM_for_Enhanced_URL_Phishing_Detection_A_Performance_Assessment
480. Introducing explainable supervised machine learning into ... <https://arxiv.org/html/2202.03212v2>
481. TabSTAR: A Foundation Tabular Model With Semantically ... <https://arxiv.org/html/2505.18125v1>
482. Machine Learning and Multi Classifier Fusion Approa <https://arxiv.org/pdf/2501.01067>

- 483. Zero-shot Meta-learning for Tabular Prediction Tasks with ... <https://arxiv.org/pdf/2502.04573>
- 484. Overtuning in Hyperparameter Optimization <https://arxiv.org/pdf/2506.19540>
- 485. TabDPT: Scaling Tabular Foundation Models <https://arxiv.org/html/2410.18164v1>
- 486. arXiv:2202.03212v2 [cs.LG] 18 Feb 2022 <https://arxiv.org/pdf/2202.03212>
- 487. Fairness Implications of Encoding Protected Categorical ... <https://arxiv.org/pdf/2201.11358>
- 488. Why doesn't CatBoost Encoding cause target leakage? <https://stats.stackexchange.com/questions/614703/why-doesnt-catboost-encoding-cause-target-leakage>
- 489. Stop the spill: The blueprint for eradicating data leakage <https://medium.com/data-science-at-microsoft/stop-the-spill-the-blueprint-for-eradicating-data-leakage-6f924e543a95>
- 490. CatBoost: A Deeper Dive <https://www.kaggle.com/code/abhinand05/catboost-a-deeper-dive>
- 491. VIIRS Nightfire https://eogdata.mines.edu/download_viirs_fire.html
- 492. Detecting forest fires with satellites (MODIS and VIIRS) <https://un-spider.org/news-and-events/news/detecting-forest-fires-satellites-modis-and-viirs>